

RetailBench: Evaluating Long-Horizon Autonomous Decision-Making and Strategy Stability of LLM Agents in Realistic Retail Environments

Anonymous ACL submission

Abstract

Large Language Model (LLM)-based agents have achieved notable success on short-horizon and highly structured tasks, yet their ability to maintain coherent decision-making over long horizons in realistic and dynamic environments remains an open challenge. We introduce *RetailBench*, a simulation benchmark designed to evaluate long-horizon autonomous decision-making in commercial scenarios, where agents must operate under stochastic demand and evolving external conditions.

Through experiments on eight state-of-the-art LLMs across progressively challenging environments, we identify and quantify four systematic failure patterns: (1) non-scalable decision-making, where agents fail to expand their decision scope as the environment grows; (2) incomplete information coverage, where critical signals are systematically ignored; (3) temporal instability, where strategies oscillate excessively across days; and (4) invalid actions, where hallucinations and constraint violations lead to system collapse. We further propose the *Evolving Strategy & Execution* framework as a reference implementation that separates high-level strategy formulation from low-level execution. In our evaluated settings with three runs per model, this framework achieves higher mean operational stability compared to day-level Reflection baselines. However, substantial performance gaps remain relative to heuristic policies, and further investigation is needed to establish statistical significance.

1 Introduction

Recent large language models (LLMs), particularly when augmented with reasoning and tool-use capabilities, have demonstrated strong performance on a variety of cognitively demanding tasks, including code editing, mathematical problem solving, and complex information retrieval (??). However, accumulating empirical evidence indicates

that such capabilities fail to generalize into robust, domain-agnostic autonomy, especially in realistic settings requiring long-horizon planning, persistent objective alignment, and stable behavioral consistency—core prerequisites for participation in real-world economic systems (??). Correspondingly, existing agent benchmarks—spanning web interaction (?????) primarily focus on short-horizon or highly structured tasks, limiting their ability to evaluate sustained interaction with complex, evolving environments. Recent studies on long-horizon autonomy consistently demonstrate that even state-of-the-art agents struggle to maintain coherent strategies over extended time spans (??).

To systematically study this challenge, we introduce *RetailBench*, a benchmark that incorporates elements from retail operations and economic modeling principles. RetailBench centers on a supermarket operation scenario that demands long-horizon decision-making, sustained interaction with a dynamic environment, and the integration of heterogeneous historical information. The benchmark evaluates whether LLM-based agents can sustain business operations under stochastic, multi-factor conditions.

We evaluate eight state-of-the-art LLMs in this environment. Our results show that current models struggle to maintain stable decision quality as the decision space expands and often fail to incorporate all relevant information. Moreover, hallucinations and economically irrational behaviors frequently emerge during long-horizon execution, leading to environment collapse and preventing sustained autonomous operation.

Our contributions are summarized as follows:

- We introduce *RetailBench*, an open-source benchmark for evaluating long-horizon decision-making in simulated retail environments, characterized by coupled pricing-inventory control, stochastic demand, and multi-day planning hori-

RetailBench Environment




Figure 1: Overview of the hierarchical supermarket environment, illustrating intra-day agent–environment interactions and end-of-day state transition dynamics.

- zons.
- Through experiments across eight models and three environment configurations, we identify and quantify four systematic failure patterns that prevent current LLM-based agents from sustaining long-horizon operations: (1) non-scalable decision-making, (2) incomplete information coverage, (3) temporal instability, and (4) invalid actions.
 - We propose the *Evolving Strategy & Execution* framework as a reference implementation that separates high-level strategy formulation from low-level execution. In our experiments, this framework achieves improved operational stability compared to day-level Reflection baselines, though substantial gaps remain relative to heuristic policies.

2 Environment Construction

2.1 Problem Formulation and Overview

We model supermarket operations as a Markov Decision Process (MDP), in which an autonomous agent manages a single retail store over a finite horizon of days. At each day t , the agent makes a sequence of operational decisions that jointly determine the store’s daily outcomes.

Design Rationale and Data Sources. The environment is designed to capture key constraints of retail operations while remaining tractable for simulation-based evaluation. Product-level demand distributions are derived from the Dominick’s Finer Foods dataset (?), which contains retail scanner data from Chicago-area stores over several years. Supply-chain parameters, including supplier price-quality relationships and delivery lead times, are constructed to reflect empirical patterns observed in operations research literature (?). External news events are synthesized using a template-based generation process that mimics financial news structures (?). However, we note that the environment remains a simplified abstraction: it does not capture competitive market dynamics, multi-store coordination, or the full complexity of real-world retail operations.

Formally, the MDP is defined as $(\mathcal{S}, \mathcal{A}, \mathcal{T}, R, \gamma)$, where $\mathcal{S}$ denotes the state space, $\mathcal{A}$ the action space, $\mathcal{T}$ the stochastic transition dynamics, R the reward function, and $\gamma \in (0, 1]$ the discount factor.

At the beginning of each day t , the agent observes the initial state $S_{t,0}$ and executes a sequence of intra-day actions indexed by k . Each action $a_{t,k} \in \mathcal{A}$ induces a transition to the next intra-day state $S_{t,k+1}$. When the agent issues the *end-today* action, the environment transitions to the initial state of the next day, $S_{t+1,0}$, according to the tran-

sition dynamics $\mathcal{T}$. Figure 1 illustrates the overall interaction process.

The environment supports long-horizon operation over more than one thousand simulated days, with episodes terminating if the store fails to pay rent for five consecutive days. In our experiments, we observe episodes ranging from 20 to 180 days depending on model performance and environment difficulty, providing sufficient temporal scale to evaluate long-horizon decision-making while remaining computationally tractable.

2.2 State Space and Observations

The environment state $S_{t,k}$ summarizes the complete operational context of the store at step k of day t . However, the agent does not observe the full state directly; instead, it receives partial observations through tool queries and action feedback. Formally, this makes the problem a Partially Observable Markov Decision Process (POMDP), though we maintain the MDP notation for simplicity and because information can be recovered through querying tools.

The complete state is composed of multiple interdependent components: $S_{t,k} = (S_{t,k}^{\text{prod}}, S_{t,k}^{\text{inv}}, S_{t,k}^{\text{sup}}, S_{t,k}^{\text{dem}}, S_{t,k}^{\text{ext}}, S_{t,k}^{\text{fin}})$.

- $S_{t,k}^{\text{prod}}$ and $S_{t,k}^{\text{inv}}$ encode product-level attributes and on-hand inventory status, including prices, shelf life, and historical sales records. Product demand is grounded in real-world retail data derived from the Dominick’s dataset (?).
- $S_{t,k}^{\text{sup}}$ represents the supply chain state, including supplier prices, quality levels, and delivery lead times, constructed to reflect empirically observed price–quality relationships (?).
- $S_{t,k}^{\text{dem}}$ captures demand-side signals such as recent customer traffic and aggregated review statistics, which influence consumer purchasing behavior (?).
- $S_{t,k}^{\text{ext}}$ represents external contextual information, including active news events with market-, category-, or product-level impact scope, synthesized to resemble real-world financial news (?).
- $S_{t,k}^{\text{fin}}$ denotes the financial state of the store, including available cash and estimated net worth, accounting for inventory depreciation over product shelf life.

Together, these components define a richly structured yet partially stochastic state space.

Observations and Tool Access. The agent receives partial observations of the state through a set of query tools. Available tools include: `query_inventory` (current stock levels), `query_sales_history` (past sales records), `query_reviews` (customer feedback), `query_supplier_info` (prices and lead times), `query_news` (current events), and `query_financial_state` (cash on hand and net worth). The agent must actively query these tools to gather information; no state information is provided automatically. This design reflects realistic operational settings where decision-makers must actively seek out relevant data. Further details of the environment design and tool specifications are deferred to Appendix ??.

2.3 Action Space

At each time step t , the agent selects an action $a_t \in \mathcal{A}$, defined as a tuple of decision components: $a_t = (a_t^{\text{price}}, a_t^{\text{repl}}, a_t^{\text{info}}, a_t^{\text{mem}}, a_t^{\text{end}})$. Each component corresponds to a distinct class of operational decisions:

- a_t^{price} denotes pricing decisions for individual SKUs;
- a_t^{repl} denotes inventory replenishment decisions, including supplier selection and order quantities;
- a_t^{info} denotes information acquisition actions, such as querying sales history, customer reviews, supplier conditions, or current news;
- a_t^{mem} denotes memory operations that allow the agent to write and retrieve persistent notes across days;
- a_t^{end} denotes a day-termination action that concludes the current decision phase.

Within a single day, the agent may execute multiple information queries and operational adjustments before issuing the a_t^{end} action to advance the environment to the next day.

2.4 Day Transition Dynamics

After the agent issues the day-termination action, the environment transitions to the next day according to $S_{t,k+1} \sim \mathcal{T}(\cdot | S_{t,k}, a_t)$, where the transition

function $\mathcal{T}$ factorizes into a sequence of stochastic updates that jointly model daily supermarket operations.

Specifically, each day-level transition consists of the following steps:

1. **Customer traffic sampling.** The total customer traffic for the day is sampled according to stochastic demand dynamics.
2. **Sales realization.** Given the sampled customer traffic and exogenous factors (e.g., promotions, seasonal effects, and news signals), the realized sales volume of each product is determined.
3. **Reviews and returns generation.** Customer reviews and product return events are generated conditional on sales outcomes, product quality, and external influences.
4. **Inventory update.** Sold units are deducted from on-hand inventory, and replenishment orders scheduled to arrive on the current day are added to stock.
5. **Financial and exogenous state update.** The agent’s financial state is updated based on realized revenues and costs, after which new exogenous information (e.g., news or market signals) for the next day is generated.

3 Evolving Strategy & Execution Framework

3.1 Framework Detail

Prior agent frameworks, including ReAct, Reflection, and Plan-and-Act (??), either do not maintain a persistent global strategy, leading to inconsistent behaviors, or revise strategies during action execution, which can induce oscillation and gradual goal drift in long-horizon environments.

To address these limitations, we propose an explicit two-stage interaction framework, termed *Evolving Strategy & Execution*, that separates strategic deliberation from operational execution. The framework maintains a persistent global strategy and restricts strategy updates to a day-level granularity to prevent excessive short-term fluctuations.

In the *Evolving Strategy* stage, the agent may invoke observation and analysis tools to examine environmental feedback and historical outcomes, but cannot execute actions that directly modify the

environment. Based on this information, it determines whether to revise the inherited strategy by adding, refining, or removing components. This stage ensures that long-term intent and planning logic remain explicit rather than implicitly entangled with immediate action choices.

In the *Execution* stage, the strategy is fixed and treated as immutable. The agent generates concrete actions strictly consistent with the current strategy, without further modification. This controlled separation enables clearer attribution between strategy and behavioral outcomes.

By alternating between these stages, the framework enforces a principled decomposition of decision-making, reduces uncontrolled strategy drift, and promotes behavioral stability and interpretability over extended horizons. An illustration is provided in Appendix ??.

3.2 Hierarchical Policy Representation

To support structured, interpretable, and temporally extended decision-making under the proposed framework, we represent the agent policy using a hierarchical abstraction that separates strategic intent from executable actions. Each policy consists of three conceptual layers:

1. **Macro Strategy**, which captures high-level managerial objectives that persist across multiple decision steps;
2. **Execution Strategy**, which encodes structured operational guidance in a machine-readable intermediate representation;
3. **Daily Actions**, which specify concrete executable operations issued to the environment.

Detailed policy configurations and example policies are provided in Appendix ??.

4 Experiment Settings

We conduct experiments under three environment configurations with increasing levels of difficulty. These configurations vary in market complexity, budget constraints, and the presence of exogenous dynamics.

4.1 Environment Configurations

We employ a heuristic policy with full access to the environment’s internal state as a calibration baseline for each environment variant. Environment parameters are tuned such that the heuristic policy remains stable across different difficulty levels while

Model	Avg. Days $\uparrow$	Avg. Daily Sales $\uparrow$	Avg. Daily Income $\uparrow$	Expiry Ratio $\downarrow$	Return Ratio $\downarrow$	Max Days $\uparrow$
<i>Framework: Evolving Strategy & Execution</i>						
GLM-4.6	52.40	174.34	124.67	0.0773	0.1293	58
Kimi-K2 (Thinking)	54.25	260.68	168.72	0.0239	0.1179	58
GPT-5.2	81.00	457.21	358.27	0.0660	0.1141	81
Average (3 models)	62.55	297.41	217.22	0.0557	0.1204	65.67
<i>Framework: Reflection (Day-Level)</i>						
GLM-4.6	55.00	160.70	125.67	0.0194	0.1176	62
Kimi-K2 (Thinking)	58.33	216.51	184.01	0.0964	0.1255	71
GPT-5.2	64.00	283.88	260.88	0.1774	0.0887	64
Average (3 models)	59.11	220.36	190.19	0.0977	0.1106	65.67
<i>Framework: Reflection (Step-Level)</i>						
GLM-4.6	51.67	92.35	77.18	0.0148	0.1072	57
Kimi-K2 (Thinking)	51.67	181.19	111.29	0.0353	0.1362	53
GPT-5.2	56.00	398.71	324.01	0.1536	0.1048	56
Average (3 models)	53.11	224.08	170.83	0.0679	0.1161	55.33
<i>Framework: Plan-and-Act</i>						
GLM-4.6	48.33	231.35	113.01	0.0198	0.1096	55
Kimi-K2 (Thinking)	48.67	170.26	105.36	0.0000	0.1056	51
GPT-5.2	64.00	323.88	193.02	0.0152	0.1096	64
Average (3 models)	53.67	241.83	137.13	0.0117	0.1083	56.67
<i>Heuristic Policy (Upper Bound, Easy)</i>						
Hand-crafted Policy	180.00	674.18	729.46	0.0266	0.0070	180

Table 1: Performance comparison of three representative large language models under four agent frameworks in the EASY environment. A hand-crafted heuristic policy is included as an approximate upper bound.

still experiencing meaningful operational pressure in Appendix ???. To evaluate models’ information-processing and external perception capabilities, we design three environment configurations:

- *Easy*: A controlled environment without dynamic news events or adaptive supplier price–quality relationships. The market contains five product categories. The agent is initialized with a budget of 10,000 and incurs a fixed daily rent of 250.
- *Middle*: A moderately complex environment that expands the product space to all twenty categories while still excluding dynamic news events and supplier adaptations. The initial budget is increased to 50,000, with a daily rent of 1,000.
- *Hard*: The most challenging and realistic environment, incorporating dynamically generated news events and time-varying supplier price–quality relationships. The market includes all twenty product categories. The agent starts with a budget of 50,000, pays a daily rent of 1,000, and receives twenty news items per day.

Detailed specifications of all environment configurations are provided in Appendix ??.

4.2 Evaluation Metrics

Metrics. We evaluate store-level operational performance using the following metrics ($\uparrow$ indicates higher is better; $\downarrow$ indicates lower is better):

- *Days* ($\uparrow$): the number of operating days before episode termination;
- *MaxDays* ($\uparrow$): the maximum number of operating days achieved across three rollouts;
- *Avg. Daily Sales* ($\uparrow$): the average number of items sold per day;
- *Avg. Daily Income* ($\uparrow$): the average money earned per day;
- *Expiry Ratio* ($\downarrow$): the fraction of products that expire before being sold;
- *Return Ratio* ($\downarrow$): the fraction of sold products that are returned by customers.

Primary Evaluation Objective. There is no single canonical score that captures all aspects of retail performance. The *Days* metric measures survival capability, *Avg. Daily Income* measures profitability, and *Expiry/Return Ratios* measure operational

efficiency. These metrics can trade off against each other: for example, reducing prices may increase sales volume but decrease margins, while holding excess inventory reduces expiry risk but increases holding costs. We report all six metrics to provide a comprehensive view, with *Days* and *Avg. Daily Income* serving as primary indicators of overall performance.

All reported metrics are averaged over three independent rollouts, each subject to a fixed maximum execution horizon.

4.3 Experimental Setup

Agent Frameworks. Preliminary experiments indicate that simple ReAct-style interaction frameworks are unstable in long-horizon settings, frequently exhibiting premature episode termination during mid-horizon execution. This observation motivates the evaluation of agent frameworks that explicitly support sustained strategic control.

We compare our proposed *Evolving Strategy & Execution* framework with step-level Reflection, day-level Reflection, and the Plan-and-Act baseline (??). For the Reflection framework, we implement two variants with different reflection frequencies: (1) step-level reflection, where the agent reflects after each action, and (2) day-level reflection, where reflection is performed once per day after completing daily actions.

Full prompt specifications for all frameworks are provided in Appendix ??.

Models and Evaluation Protocol. We evaluate a diverse set of contemporary large language models, including Qwen-235B (Thinking) (?), Kimi K2 (Thinking) (?), GLM-4.6 (?), DeepSeek-V3.2-Exp (?), Gemini-3-Flash-Preview (?), Grok-4.1 Fast (?), GPT-5-Mini (?), and GPT-5.2 (?). To manage the substantial token costs of long-horizon rollouts, lower-cost variants are used for closed-source models when available.

Each model is evaluated across three environment configurations, with three independent rollouts per configuration; results are averaged across runs. To isolate the effect of the agent framework, we conduct controlled comparisons between our framework and baseline methods in the *Easy* environment under identical conditions. We further implement a hand-crafted policy based on privileged internal knowledge unavailable to the agents, which serves as an approximate upper bound on performance.

5 Results

5.1 Performance Comparison between Agent Frameworks

Due to computational budget constraints, we conduct a comprehensive four-framework comparison on three representative models: GLM-4.6, Kimi-K2 (Thinking), and GPT-5.2¹. As reported in Table ??, *Evolving Strategy & Execution* achieves higher mean sales and profit, and lower mean product expiration rates, compared to alternative agent frameworks in our evaluated settings.

We attribute this improvement to two design features. First, restricting strategy updates to day-level granularity reduces oscillatory behavior compared to step-level reflection, where strategies may change after every action. Second, explicit strategy representation improves consistency across decisions within each day, as evidenced by lower execution strategy churn rates compared to baselines. However, we observe that even our framework exhibits substantial day-to-day strategy variability, suggesting that maintaining stable long-horizon strategies remains a fundamental challenge for current LLMs.

Based on these results, we select the two strongest-performing frameworks for broader evaluation across all eight models. When averaged over all models, our proposed framework demonstrates superior mean performance metrics relative to *Reflection (Day-Level)*.

Despite these gains, a substantial gap remains relative to the hand-crafted heuristic policy, which serves as an approximate upper bound. This gap highlights the current limitations of LLM-based agents in long-horizon decision-making. Notably, Table ?? also suggests that larger closed-source models, such as GPT-5.2, tend to achieve more stable and sustained store operation compared to smaller or open-source counterparts, indicating that model capacity remains an important factor alongside framework design.

Table ?? presents the complete results across all eight evaluated models. Across models, we observe recurring and systematic failure modes in the current environment, suggesting that the perfor-

¹Due to computational constraints, our current evaluation uses a limited number of experimental runs. The reported performance differences should be interpreted as preliminary observations of mean performance rather than statistically significant conclusions. We plan to extend our evaluation with paired-seed experimental designs and formal significance testing in future work.




Figure 2: Category-level sales and profit per category across Easy, Middle, and Hard environments. Results are shown for three representative models.

mance bottlenecks are structural rather than model-specific.

5.2 Performance across Environments with Varying Difficulty

As environment difficulty increases, all models exhibit consistent performance degradation, including shorter operational durations, higher expiry and return ratios, and reduced long-horizon stability. Although the expansion of product categories in the Middle and Hard settings enables higher aggregate sales and profit for some models, *Sales per Category* and *Profit per Category* decline substantially (Figure ??), indicating persistent challenges in effective resource allocation within increasingly high-dimensional decision spaces.

The relatively small performance gap between the Middle and Hard environments can be partly attributed to the delayed impact of intensified news dynamics, whose effects typically unfold over longer time horizons. Appendix ?? reports the full results across all environments under the proposed framework. Overall, while models demonstrate limited adaptability as task complexity increases in our experiments, their performance remains substantially below the heuristic upper bound, suggesting challenges that current LLMs face in long-horizon decision-making within the RetailBench environment. Whether these patterns generalize to other domains remains an open question.

6 Analysis

Based on both quantitative evaluation results and manual inspection of trajectories, we identify several key factors that explain why current models fail to operate reliably in the environment.

We analyze these failure modes from four complementary perspectives.

Model	Context	Easy		Middle		Hard	
		SKU ↑	Cat ↑	SKU ↑	Cat ↑	SKU ↑	Cat ↑
DeepSeek-V3.2 (Exp.)	128K	4.75	3.31	6.83	<u>5.34</u>	6.64	<u>5.45</u>
Gemini-3 (Fast)	1M	8.72	4.34	13.60	12.43	21.57	13.56
GLM-4.6	200K	4.82	2.98	3.23	2.46	<u>7.08</u>	5.22
OpenAI-5.1 Mini	400K	3.66	1.94	4.10	4.10	6.67	4.79
Grok-4.1 Fast	200K	<u>5.78</u>	<u>3.68</u>	4.87	3.08	5.40	3.20
Kimi-K2 (Thinking)	256K	5.06	3.09	5.17	4.77	6.91	4.75
Qwen-235B (Thinking)	256K	4.59	3.67	3.17	2.36	5.11	4.51
Heuristic Strategy	–	9.03	4.87	35.34	18.61	34.82	18.52

Table 2: SKU and Category counts sold by different models across difficulties. Context denotes the maximum supported context length of each model. **Bolded** indicates the best model; underlined indicates the second best (Human excluded).

Model	Supplier	Inventory	Return	Rating	Price	Review	Sales	History
DeepSeek-V3.2 (Exp.)	41.4	100.0	26.7	75.4	6.6	0.6	89.2	0.0
Gemini-3 (Fast)	69.6	100.0	41.0	82.3	67.7	6.3	87.6	2.9
GLM-4.6	93.9	98.3	79.8	77.8	84.1	5.1	96.3	0.0
OpenAI-5.1 Mini	58.8	99.0	5.1	78.6	3.9	10.6	84.7	0.3
Grok-4.1 Fast	89.9	99.8	84.0	94.0	88.6	36.0	96.4	0.3
Kimi-K2 (Thinking)	57.2	90.4	25.9	51.4	36.2	11.0	64.0	0.5
Qwen-235B (Thinking)	76.8	78.3	31.5	58.2	19.8	23.6	74.6	0.0
Average	462	69.7	95.1	42.0	74.0	43.8	84.7	1.0

Table 3: Percentage of days on which each model queries specific information sources when making decisions for SKUs included in the final daily strategy. Higher values indicate greater reliance on the corresponding data source.

6.1 Non-scalable Decision-Making Capability

As shown in Appendix ??, models achieve reasonable performance in the Easy setting but exhibit consistent degradation as the environment scales in complexity. To better understand this phenomenon, we analyze both the average number of stock keeping units (SKUs) and product categories sold per day, as well as the number of SKUs and categories explicitly considered during the *Evolving Strategy* phase (Tables ?? and Appendix ??).

Across most models, decision-making capability does not scale proportionally with the size of the environment. Instead, performance remains relatively flat despite a substantial increase in the number of available options. Notably, Gemini-3 (Fast), which supports the largest maximum context window, performs better than other models in this respect, suggesting that larger context capacity helps retain salient information even when the effective interaction context is constrained.

Nevertheless, even the strongest models fail to cover the full decision space. This indicates that current systems are unable to expand their effective decision-making scope as the environment grows, leading to systematic performance degradation in larger and more complex settings.

Model	Macro Strategy			Execution Strategy		
	Std_diff ($\downarrow$)	MAC ($\downarrow$)	TV ($\downarrow$)	Std_diff ($\downarrow$)	MAC ($\downarrow$)	TV ($\downarrow$)
DeepSeek-V3.2(Exp.)	0.130	0.090	5.00	0.289	0.144	7.93
Gemini-3 (Fast)	0.136	0.085	3.82	0.293	0.225	10.18
GLM-4.6	0.131	0.096	4.52	0.264	0.209	9.87
OpenAI-5.1 Mini	0.069	0.051	2.50	0.229	0.166	8.00
Grok-4.1 Fast	0.079	0.045	2.71	0.204	0.151	9.39
Kimi K2(Thinking)	0.179	0.133	6.95	0.335	0.271	14.08
Qwen-235B (Thinking)	0.240	0.189	6.51	0.391	0.306	10.57

Table 4: Temporal instability metrics of macro- and execution-level strategy similarity in the Easy environment. All metrics are lower-is-better ($\downarrow$). Larger values indicate greater temporal instability. Column-wise maximum values are highlighted in bold.

6.2 Incomplete Decision-Making Due to Limited Information Coverage

We analyze the information sources that models attend to when making operational decisions by examining the data queried for the set of SKUs included in each day’s final strategy. This analysis reveals a clear concentration of attention on a limited subset of signals.

Specifically, most models primarily rely on supplier prices, inventory levels, SKU ratings, and historical sales records when deciding how to operate selected SKUs in Table ?? . In contrast, several other critical signals—such as recent customer reviews, return rates, and current selling prices—are consistently underutilized or entirely ignored.

Further correlation analysis shows a strong positive relationship between the frequency of SKU reviews queries and average daily sales performance in Appendix ?? . This observation aligns with the underlying environment dynamics and indicates that incomplete information coverage is a key factor limiting decision quality. Overall, these results suggest that models often fail to perform sufficiently comprehensive information gathering, leading to systematically suboptimal operational decisions.

6.3 Temporal Instability in Execution-Level Decision-Making

Beyond limitations in decision scalability and information coverage, we identify temporal instability in execution-level decision-making as a key contributor to long-horizon failure. Even under relatively stable environmental conditions, agents frequently revise their strategies across consecutive days, leading to inconsistent execution trajectories.

Measuring Strategy Similarity. We quantify temporal instability by measuring the similarity between strategies on adjacent days at both macro

and execution levels. Macro strategy similarity is assessed using an LLM-based prompt that evaluates semantic consistency between consecutive high-level plans. Execution strategy similarity is computed via set-based Jaccard similarity over key fields, including `focus_skus`, `sku_supplier_mapping`, `news_to_monitor`, and `sku_to_monitor`, with the final execution similarity obtained by averaging across fields.

Instability Metrics. To characterize temporal fluctuations, we compute three complementary metrics: the standard deviation of first-order differences (*Std_diff*) to capture short-term volatility, the mean absolute change (*MAC*) to estimate typical day-to-day variation, and total variation (*TV*) to reflect cumulative long-term instability.

Across all three environment configurations, we observe consistent temporal instability in both macro- and execution-level strategies. For clarity, we present detailed results from the Easy environment in Table ?? and Appendix ?? . Despite its reduced complexity, the Easy setting already exhibits pronounced temporal fluctuations. In particular, macro strategies remain relatively stable over time, whereas execution strategies show substantially larger variability across most models. This effect is especially pronounced for Qwen-235B (Thinking), which also performs poorly across all environments.

Overall, these results indicate that long-horizon failures arise not only from suboptimal strategy formulation, but more fundamentally from the inability to maintain temporally consistent execution policies over extended horizons.

6.4 Hallucinations and Invalid Actions

Finally, manual inspection reveals recurrent failure patterns that directly break planning correctness and action validity. We distinguish two closely related but practically different issues:

Hallucinations in reasoning Models occasionally generate reasoning traces that reference non-existent SKUs or fabricate numerical quantities, and subsequently incorporate these hallucinated elements into multi-step plans (examples are provided in Appendix ??). As a result, decisions become misaligned with the true environment state, even when the overall planning structure appears internally coherent.

Invalid or irrational actions. Models sometimes output actions that violate basic constraints or are inconsistent with historical demand, such as negative order quantities or implausible pricing in Appendix ???. Even though state-modifying operations are restricted to only a small set of tools, these invalid actions occur with non-negligible frequency and can quickly destabilize the system in long-horizon operation.

In realistic operational settings, both hallucinations and invalid actions would be unacceptable and could directly trigger system collapse.

7 Related Work

Long-horizon planning benchmarks for LLMs. In parallel, a growing body of benchmarks targets long-horizon planning abilities of large language models across structured and interactive environments. PlanBench focuses on classical plan generation, while WebShop, Mind2Web, and ScienceWorld (???) evaluate multi-step interaction, error recovery, and adaptive behavior in dynamic settings. More recent benchmarks—such as HeroBench, OdysseyBench, UltraHorizon (???), and explicitly stress extended, interdependent decision sequences that require persistent memory, hierarchical reasoning, and long-term strategy maintenance. Collectively, these benchmarks suggest that while short-horizon tasks are increasingly tractable, robust long-horizon planning and execution remain a central challenge for LLM-based agents.

Long-horizon agent frameworks. Recent advances in agent design address these challenges by introducing structured frameworks that decouple high-level planning from low-level execution. Approaches such as Plan-and-Act (?) and EA-GLT (?) adopt planner-executor architectures to support hierarchical decision-making, dynamic replanning, and improved execution stability over long horizons. Extensions to multi-agent settings ELHPlan (?) and plan-aware context management frameworks PAACE (?) further emphasize task decomposition, explicit strategy representation, and memory-aware context control. Together, these works indicate that scalable long-horizon autonomy increasingly relies on structured planning modules and controlled execution mechanisms rather than monolithic end-to-end policies.

Classical operations research baselines. A significant limitation of our current evaluation is the

Table 5: Comparison of RetailBench with related long-horizon benchmarks.

Benchmark	Domain	Horizon	Economic Obj.	Op
WebArena (?)	Web tasks	Short-Horizon	No	
Mind2Web (?)	Web tasks	Medium-Horizon	No	
VendingBench (?)	Vending	Long-Horizon	Limited	
OdysseyBench (?)	Games	Long-Horizon	No	
RetailBench	Retail Ops.	Long-Horizon	Yes	

absence of strong classical operations research and control theory baselines. Retail operations have been extensively studied in the OR literature, with well-established solutions including base-stock policies for inventory control (?), newsvendor models for single-period decisions, and dynamic pricing algorithms (?). Methods such as Model Predictive Control (MPC) and approximate dynamic programming have also been applied to multi-period inventory and pricing problems. Our current evaluation compares only against prompting-based LLM agent frameworks; we do not include these classical OR baselines, which likely represent strong oracles for many subtasks. The relative performance of LLM-based agents compared to these well-established methods remains an important open question.

Positioning relative to existing benchmarks. RetailBench complements existing benchmarks by targeting a specific combination of properties not jointly addressed in prior work: (1) long-horizon planning with episodes up to 180 days (limited by rent defaults), (2) economic optimization objectives (profit maximization under inventory and pricing constraints), (3) coupled decision spaces where pricing affects inventory turnover and vice versa, (4) stochastic dynamics from demand uncertainty, news shocks, and supply-chain delays, and (5) open-source availability for reproducible evaluation. While benchmarks such as WebArena and Mind2Web evaluate multi-step tool use, and VendingBench and OdysseyBench assess long-horizon planning, RetailBench is distinguished by its focus on economically grounded retail operations with realistic temporal dependencies and delayed rewards. Table ?? provides a detailed comparison.

8 Conclusion

This paper introduces RetailBench, a benchmark for evaluating long-horizon decision-making in simulated retail environments. RetailBench models supermarket operations as a stochastic, multi-factor,

and temporally extended process, requiring agents to reason jointly about pricing, inventory, information acquisition, and financial sustainability.

Our primary contribution is the identification and quantification of four systematic failure patterns that prevent current LLM-based agents from sustaining long-horizon operations: non-scalable decision-making, incomplete information coverage, temporal instability, and invalid actions. We further propose the Evolving Strategy & Execution framework as a reference implementation for benchmark participants. In our evaluated settings with three runs per model, this framework achieves higher mean operational stability compared to day-level Reflection baselines. However, substantial performance gaps remain relative to heuristic policies, and observed performance variations across random seeds indicate that further investigation is needed to establish statistical significance. Performance also degrades sharply as environment complexity increases. RetailBench provides a testbed for advancing research on long-horizon decision-making and agentic reasoning, though external validity beyond the simulated retail setting remains to be established.

9 Limitations

Limitations

Our work has several important limitations that constrain the interpretation and generalizability of our results.

Evaluation scope. Our evaluation is conducted in a simulated single-store supermarket environment. While the environment incorporates realistic elements such as stochastic demand, perishable inventory, and exogenous shocks, it remains a simplified abstraction of real-world retail operations. We do not model multi-store coordination, competitive market dynamics, or strategic interactions among multiple autonomous agents. External validity beyond this simulated setting remains to be established.

Statistical rigor. Due to computational constraints, our current evaluation uses a limited number of experimental runs (three runs per model in most settings). The reported performance differences should be interpreted as preliminary observations of mean performance rather than statistically significant conclusions. Rank orderings may be sensitive to random seed selection, and effect sizes may vary with additional repetitions. We plan to

extend our evaluation with paired-seed experimental designs and formal significance testing in future work.

Causal claims. Our comparison of agent frameworks demonstrates correlations between framework design and performance metrics, but does not establish causal relationships. Performance differences may be attributable to factors including token budget, prompt engineering, or model-specific behaviors rather than framework architecture alone. Controlled ablation studies are needed to isolate the causal impact of individual framework components.

Baseline strength. Our current baselines include rule-based heuristics and prompting-based agent framework variants, but do not include strong classical operations research methods such as (s, S) inventory policies, newsvendor-based ordering strategies, or dynamic pricing algorithms. Classical OR methods such as Model Predictive Control and approximate dynamic programming have been extensively studied for retail operations and likely represent strong oracles for many subtasks. The relative performance of LLM-based agents compared to these well-established methods remains unknown and is an important direction for future investigation.

Learning and adaptation. Our evaluation focuses on prompting-based LLM agents without parameter updates or inter-episode learning. Stronger performance may be achievable through reinforcement learning, fine-tuning, or hybrid neuro-symbolic approaches. Our findings should be interpreted as characterizing zero-shot prompting performance rather than the full potential of learned agent systems.

Failure mode analysis. While we identify recurring patterns such as hallucinations and economically irrational actions, we do not propose explicit algorithmic mechanisms to enforce economic rationality or factual grounding during execution. Addressing these failure modes through constraint-aware action control remains an important open problem.

Domain specificity. Our findings are specific to the retail operations domain and the RetailBench environment. Transferability to other decision-making domains—such as manufacturing, logistics, or financial trading—requires empirical validation and should not be assumed without evidence.

References

- Dario Amodei. 2024. Machines of loving grace. <https://www.darioamodei.com/essay/machines-of-loving-grace>. Accessed: 2025-12-01.
- Andon Labs. 2025. Vending-bench 2: A benchmark for long-horizon business simulation. <https://andonlabs.com/evals/vending-bench-2>. Accessed: 2025-12-10.
- Petr Anokhin, Roman Khalikov, Stefan Rebrikov, Viktor Volkov, Artyom Sorokin, and Vincent Bissonnette. 2025. Herobench: A benchmark for long-horizon planning and structured reasoning in virtual worlds. *Preprint*, arXiv:2508.12782.
- ashraq. 2025. financial-news-articles. <https://huggingface.co/datasets/ashraq/financial-news-articles>. Accessed: 2025-12-01.
- Axel Backlund and Lukas Petersson. 2025. Vending-bench: A benchmark for long-term coherence of autonomous agents. *Preprint*, arXiv:2502.15840.
- DeepSeek-AI, Aixin Liu, Aoxue Mei, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, Chen Dong, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenhao Xu, Chong Ruan, Damai Dai, Daya Guo, Dejian Yang, and 245 others. 2025. Deepseek-v3.2: Pushing the frontier of open large language models. *Preprint*, arXiv:2512.02556.
- Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Samuel Stevens, Boshi Wang, Huan Sun, and Yu Su. 2023. Mind2web: Towards a generalist agent for the web. *Preprint*, arXiv:2306.06070.
- Lutfi Eren Erdogan, Nicholas Lee, Sehoon Kim, Suhong Moon, Hiroki Furuta, Gopala Anumanchipalli, Kurt Keutzer, and Amir Gholami. 2025a. Plan-and-act: Improving planning of agents for long-horizon tasks. *Preprint*, arXiv:2503.09572.
- Lutfi Eren Erdogan, Nicholas Lee, Sehoon Kim, Suhong Moon, Hiroki Furuta, Gopala Anumanchipalli, Kurt Keutzer, and Amir Gholami. 2025b. Plan-and-act: Improving planning of agents for long-horizon tasks. *Preprint*, arXiv:2503.09572.
- Dave Fedewa, Chris Holder, Wynn Teichner, and Ben Wiseman. 2021. Five-star growth: Using online ratings to design better products. *McKinsey & Company*. Accessed: 2025-11-30.
- Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma, Liang Chen, Runxin Xu, Zhengyang Tang, Benyou Wang, Daoguang Zhan, Shanghaoran Quan, Ge Zhang, Lei Sha, Yichang Zhang, Xuancheng Ren, Tianyu Liu, and Baobao Chang. 2024. Omni-math: A universal olympiad level mathematic benchmark for large language models. *Preprint*, arXiv:2410.07985.
- Google DeepMind. 2024. Gemini 3 flash (preview). <https://deepmind.google/technologies/gemini/>. Large multimodal language model, preview version.
- Dhruv Grewal, Jens Nordfalt, Anne Roggeveen, Rainer Olbrich, and Hans Christian Jansen. 2014. Price-quality relationship in pricing strategies for private labels. *Journal of Product and Brand Management*, 23(6):429–438.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. 2024. Swe-bench: Can language models resolve real-world github issues? *Preprint*, arXiv:2310.06770.
- University of Chicago Booth School of Business Kilts Center for Marketing. Dominick’s dataset. <https://www.chicagobooth.edu/research/kilts/research-data/dominicks>. Accessed: 2025-11-01.
- Thomas Kwa, Ben West, Joel Becker, Amy Deng, Katharyn Garcia, Max Hasin, Sami Jawhar, Megan Kinniment, Nate Rush, Sydney Von Arx, Ryan Bloom, Thomas Bradley, Haoxing Du, Brian Goodrich, Nikola Jurkovic, Luke Harold Miles, Seraphina Nix, Tao Lin, Neev Parikh, and 6 others. 2025. Measuring ai ability to complete long tasks. *Preprint*, arXiv:2503.14499.
- Shaobin Ling, Yun Wang, Chenyou Fan, Tin Lun Lam, and Junjie Hu. 2025. Elhplan: Efficient long-horizon task planning for multi-agent collaboration. *Preprint*, arXiv:2509.24230.
- Haotian Luo, Huaisong Zhang, Xuelin Zhang, Haoyu Wang, Zeyu Qin, Wenjie Lu, Guozheng Ma, Haiying He, Yingsha Xie, Qiyang Zhou, Zixuan Hu, Hongze Mi, Yibo Wang, Naiqiang Tan, Hong Chen, Yi R. Fung, Chun Yuan, and Li Shen. 2025. Ultrahorizon: Benchmarking agent capabilities in ultra-long-horizon scenarios. *Preprint*, arXiv:2509.21766.
- Daniel McFadden. 1974. Conditional logit analysis of qualitative choice behavior. In Paul Zarembka, editor, *Frontiers in Econometrics*, pages 105–142. Academic press, New York.
- METR. 2025. Measuring ai ability to complete long tasks. METR blog.
- Grégoire Mialon, Clémentine Fourrier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. 2023. Gaia: a benchmark for general ai assistants. *Preprint*, arXiv:2311.12983.
- NoF1.ai. 2025. Alpha arena — exploring the limits of large language models as quant traders. <https://nof1.ai/blog/TechPost1>. Accessed: 2025-12-10.
- OpenAI. 2025a. Gpt-5 mini. <https://platform.openai.com/docs/models/gpt-5-mini>. Large language model — cost-efficient GPT-5 variant.

OpenAI. 2025b. Introducing gpt-5.2 . Accessed: 2026-02-19.	962
Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, Hugh Zhang, Chen Bo Calvin Zhang, Mohamed Shaaban, John Ling, Sean Shi, Michael Choi, Anish Agrawal, Arnav Chopra, Adam Khoja, Ryan Kim, Richard Ren, Jason Hausenloy, Oliver Zhang, Mantas Mazeika, and 1093 others. 2025. Humanity’s last exam . <i>Preprint</i> , arXiv:2501.14249.	963
Ibrahim Shar, Wenhuan Sun, Haiyan Wang, and Chetan Gupta. 2022. Deep reinforcement learning toward robust multi-echelon supply chain inventory optimization . pages 1385–1391.	964
Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language agents with verbal reinforcement learning . <i>Preprint</i> , arXiv:2303.11366.	965
Shuzheng Si, Haozhe Zhao, Kangyang Luo, Gang Chen, Fanchao Qi, Minjia Zhang, Baobao Chang, and Maosong Sun. 2025. A goal without a plan is just a wish: Efficient and effective global planner training for long-horizon agent tasks . <i>Preprint</i> , arXiv:2510.05608.	966
GLM Team, Aohan Zeng, Xin Lv, Qinkai Zheng, Zhenyu Hou, Bin Chen, Chengxing Xie, Cunxiang Wang, Da Yin, Hao Zeng, Jiajie Zhang, Kedong Wang, Lucen Zhong, Mingdao Liu, Rui Lu, Shulin Cao, Xiaohan Zhang, Xuancheng Huang, Yao Wei, and 152 others. 2025a. Glm-4.5: Agentic, reasoning, and coding (arc) foundation models . <i>Preprint</i> , arXiv:2508.06471.	967
Kimi Team, Yifan Bai, Yiping Bao, Guanduo Chen, Jiahao Chen, Ningxin Chen, Ruijue Chen, Yanru Chen, Yuankun Chen, Yutian Chen, Zhuofu Chen, Jialei Cui, Hao Ding, Mengnan Dong, Angang Du, Chenzhuang Du, Dikang Du, Yulun Du, Yu Fan, and 150 others. 2025b. Kimi k2: Open agentic intelligence . <i>Preprint</i> , arXiv:2507.20534.	968
Qwen Team. 2025a. Qwen3 technical report . <i>Preprint</i> , arXiv:2505.09388.	969
The Terminal-Bench Team. 2025b. Terminal-bench: A benchmark for ai agents in terminal environments .	970
Mark D. Uncles. 1987. Discrete choice analysis: Theory and application to travel demand . <i>Journal of the Operational Research Society</i> , 38(4):370–371.	971
Ruoyao Wang, Peter Jansen, Marc-Alexandre Côté, and Prithviraj Ammanabrolu. 2022. Scienceworld: Is your agent smarter than a 5th grader? <i>Preprint</i> , arXiv:2203.07540.	972
Weixuan Wang, Dongge Han, Daniel Madrigal Diaz, Jin Xu, Victor Rühle, and Saravan Rajmohan. 2025. Odysseybench: Evaluating llm agents on long-horizon complex office application workflows . <i>Preprint</i> , arXiv:2508.09124.	973
Jason Wei, Zhiqing Sun, Spencer Papay, Scott McKinney, Jeffrey Han, Isa Fulford, Hyung Won Chung, Alex Tachard Passos, William Fedus, and Amelia Glaese. 2025. Browsecomp: A simple yet challenging benchmark for browsing agents . <i>Preprint</i> , arXiv:2504.12516.	974
xAI. 2025. Grok 4.1 fast and agent tools api. https://x.ai/news/grok-4-1-fast . Accessed [Your Access Date].	975
Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. 2023a. Webshop: Towards scalable real-world web interaction with grounded language agents . <i>Preprint</i> , arXiv:2207.01206.	976
Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023b. React: Synergizing reasoning and acting in language models . <i>Preprint</i> , arXiv:2210.03629.	977
Kamer Ali Yuksel. 2025. Paace: A plan-aware automated agent context engineering framework . <i>Preprint</i> , arXiv:2512.16970.	978
Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. 2024. Webarena: A realistic web environment for building autonomous agents . <i>Preprint</i> , arXiv:2307.13854.	979
A Environment Configuration Details	980
A.1 State Space Decomposition	981
We construct the environment using real-world retail data from the Dominick’s dataset (?). From the 20 available product categories, we select 96 SKUs with the most complete and informative sales records. The richness of these observations enables reliable modeling of the relationship between pricing decisions and realized demand.	982
Key Symbols. Let $\mathcal{J}$ denote the set of selected SKUs, with cardinality $ \mathcal{J} = J$ (where $J = 96$ in our experiments). Each SKU $j \in \mathcal{J}$ belongs to a product category $\text{cat}(j) \in \{1, \dots, 20\}$. For each SKU j , let $\mathcal{K}(j)$ denote its associated supplier set, with $ \mathcal{K}(j) = 5$.	983
A.1.1 Product State S_t^{prod}	984
For each SKU j , we maintain a set of static attributes together with time-varying operational signals:	985
$S_{t,j}^{\text{prod}} = (\text{id}_j, \text{desc}_j, L_j, p_{j,t}, \mathbf{h}_{j,t}^{\text{sales}}, \mathbf{r}_{j,t}). \quad (1)$	986
Here, id_j and desc_j denote the unique identifier and textual description of SKU j , respectively. L_j denotes the shelf life (in days), and $p_{j,t}$ is the retail	987

price on day t . The term $\mathbf{h}_{j,t}^{\text{sales}}$ encodes historical sales statistics, while $\mathbf{r}_{j,t}$ summarizes aggregated customer review signals, such as average rating and review volume.

Data grounding. The SKU set and cost priors are constructed from the Dominick’s dataset (?). Historical sales records from the same source are used to calibrate the demand model.

A.1.2 Inventory State S_t^{inv}

Inventory is represented at the SKU level with age tracking. Let I_{jt} denote the on-hand units of SKU j at the beginning of day t . To support shelf-life constraints and depreciation, we track the arrival time and remaining shelf life of each unit.

Capacity constraint and pending queue. Let Cap denote the total inventory capacity of the store. If incoming replenishment orders would violate this constraint, excess units are placed into a first-in-first-out (FIFO) pending queue Q_t and become available only after sufficient inventory space is released:

$$\sum_{j \in \mathcal{J}} \sum_{a=0}^{L_j} I_{jt}^{(a)} \leq \text{Cap}. \quad (2)$$

Stockout signal. When realized demand exceeds the available sellable inventory of SKU j on day t , a stockout indicator is triggered. This signal is recorded at the end of day t and exposed to the agent as explicit operational feedback.

Expiration and destruction. At each day transition, inventory units whose residence time exceeds their shelf life are removed from the system.

A.1.3 Supply Chain State S_t^{sup}

Each SKU j is associated with a set of suppliers $k \in \mathcal{K}(j)$. The supplier state includes procurement price c_{jk} , quality level q_{jk} , and lead-time distribution:

$$S_{t,j}^{\text{sup}} = \{(c_{jk}, q_{jk}, \mathcal{L}_{jk}) : k \in \mathcal{K}(j)\}. \quad (3)$$

Lead time $\ell_{jk} \sim \mathcal{L}_{jk}$ is sampled when an order is placed. Procurement prices are discretized into tiers using Dominick’s cost signals (?), following empirically observed positive correlations between price tier and quality level (?).

Enforced diversity. To avoid degenerate supplier configurations, we enforce that each SKU has (i) one supplier with maximal quality, (ii) one

supplier with minimal price and minimal quality, and remaining suppliers spanning intermediate price-quality levels.

A.1.4 Demand Signals and Demand Generation

Demand-related components capture both observable market signals and the stochastic process governing realized demand. We define

$$S_t^{\text{dem}} = (N_t, \{\mathbf{r}_{j,t}\}_{j \in \mathcal{J}}), \quad (4)$$

where N_t denotes daily customer traffic.

Customer traffic. Daily customer traffic is derived from the Dominick’s dataset.

Consumer choice and demand generation.

Given customer traffic N_t , prices $\{p_{jt}\}_{j \in \mathcal{J}}$, review signals $\{\mathbf{r}_{j,t}\}$, and external news $\mathcal{E}_t$, we generate realized demand using a discrete-choice model. Following standard practice in empirical economics, we adopt a Multinomial Logit (MNL) framework (??).

For SKU j , the raw utility is defined as

$$U_{jt}^{\text{raw}} = \alpha_j + \beta_j p_{jt} + \varepsilon_{jt}, \quad (5)$$

where α_j captures intrinsic preference, β_j models price sensitivity, and ε_{jt} represents idiosyncratic shocks.

We further incorporate review effects, news signals, and within-category substitution:

$$\begin{aligned} \tilde{U}_{jt} &= U_{jt}^{\text{raw}} + \Delta_j(\mathbf{r}_{j,t}, \mathcal{E}_t) \\ &+ \sum_{\substack{i \neq j \\ \text{cat}(i) = \text{cat}(j)}} \gamma_{ji} \exp(\tilde{U}_{it}). \end{aligned} \quad (6)$$

With the outside option utility normalized to zero, the purchase probability for SKU j on day t is

$$p_{jt}^* = \frac{\exp(\tilde{U}_{jt})}{1 + \sum_{i=1}^J \exp(\tilde{U}_{it})}. \quad (7)$$

Realized demand. Given traffic N_t and purchase probabilities $\{p_{jt}^*\}$, potential demand is sampled as

$$y_{jt} \sim \text{Binomial}(N_t, p_{jt}^*), \quad (8)$$

and realized sales are capped by available on-hand inventory.

A.1.5 External Information S_t^{ext} (News Module)

We maintain a set of active news events $\mathcal{E}_t$. Each event $e \in \mathcal{E}_t$ is represented as

$$e = (\text{type}, \text{scope}, \text{target}, \text{side}, \text{sign}, \eta, \text{text}, \text{ttl}) \quad (9)$$

where $\text{scope} \in \{\text{macro}, \text{category}, \text{product}, \text{neutral}\}$, $\text{side} \in \{\text{demand}, \text{supply}, \text{both}\}$, $\text{sign} \in \{+1, -1\}$, η denotes impact magnitude, and ttl is the time-to-live in days. News texts are synthesized using an LLM to resemble financial news corpora (?).

Impact application. News impacts are incorporated into (i) demand utilities and/or (ii) supplier prices depending on side and scope. Neutral news has $\eta = 0$ by design.

A.1.6 Financial State S_t^{fin}

Let F_t denote available funds at the start of day t . Net worth is computed as

$$\text{NW}_t = F_t + \sum_{j \in \mathcal{J}} \sum_{a=0}^{L_j} I_{jt}^{(a)} \cdot v_j(a), \quad (10)$$

where $v_j(a)$ denotes the per-unit value at age a . We adopt linear shelf-life depreciation:

$$v_j(a) = c_j \cdot \max\left(0, 1 - \frac{a}{L_j}\right), \quad (11)$$

with c_j denoting a reference procurement cost (e.g., the mean or realized supplier cost).

A.2 Policy Detail

A.2.1 Policy Presentation

To support structured and interpretable decision making, we represent the agent policy using a hierarchical abstraction that separates strategic intent from executable actions. Specifically, each policy is composed of three layers:

1. **Macro Strategy**, which captures high-level managerial principles that persist across days;
2. **Execution Strategy**, which encodes structured operational guidance in a machine-readable form;

3. **Daily Actions**, which enumerate concrete executable operations submitted to the environment.

This design enables the agent to reason at different temporal and semantic granularities, while maintaining a clear separation between planning and execution.

Macro Strategy. The macro strategy consists of a set of natural-language statements that describe high-level objectives (e.g., prioritizing inventory turnover or focusing on high-margin products). These statements are non-executable and serve as persistent guidance for downstream decision making.

Execution Strategy. The execution strategy consists of six key components. `focus_skus` specifies the SKUs that require immediate attention. `sku_supplier_mapping` denotes the corresponding suppliers for each SKU. `news_to_monitor` identifies relevant news signals to track. `skus_to_reorder` indicates SKUs that require replenishment. `price_adjustment` specifies the SKUs whose prices should be adjusted along with the corresponding adjustment magnitudes. `sku_to_monitor` denotes SKUs under observation. The other field captures additional execution directives.

Daily Actions. Daily actions consist of two operation types: `place_order`, which places purchase orders for selected SKUs, and `modify_sku_price`, which adjusts the prices of specified SKUs.

Strategy-Execution Protocol. Policy usage follows a two-phase protocol. In the *strategy phase*, the agent analyzes the environment and constructs its macro and execution strategies. In the subsequent *execution phase*, the finalized strategy is consumed to generate executable daily actions. The strategy is immutable during execution, enforcing a clear separation between deliberation and action.

A.2.2 Policy Example

See in Table ??.

A.3 Environment Setting

A.3.1 Easy Environment Configuration

The Easy configuration is designed for simpler scenarios with limited resources and a reduced category set.

- **Time Range:**

– Data begin time: 06/06/91	1175	Frozen_Juices, Oatmeal, Paper_Towels,	1215
– Data end time: 12/31/95	1176	Snack_Crackers	1216
– Store begin time: 09/07/91	1177	Soft_Drinks, Toothpastes	1217
– Store ID: 15	1178		
• Financial Parameters:	1179	• Category Effects: All categories have a uniform effect of -0.2	1218
– Initial funds: 10,000			1219
– Daily rent: 250		• Random Seed: 42	1220
– Inventory capacity: 10,000	1181		
• Feature Enablement:	1182	A.3.3 Hard Environment Configuration	1221
– Review enabled: True	1183	The Hard configuration uses dynamic data with news events enabled, providing the most complex simulation environment.	1222
– Review ratio: 0.02	1184		1223
– News enabled: False	1185	• Time Range: Same as other configurations	1224
• Selected Categories (5 categories):	1186		1225
– Bathroom_Tissues	1187	• Financial Parameters:	1226
– Canned_Soup	1188	– Initial funds: 50,000	1227
– Cigarettes	1189	– Daily rent: 1,000	1228
– Front_end_candies	1190	– Inventory capacity: 40,000	1229
– Soft_Drinks	1191	• Feature Enablement:	1230
• Category Effects: All categories have a uniform effect of -0.2	1192	– Review enabled: True	1231
	1193	– Review ratio: 0.02	1232
• Random Seed: 42 (for reproducibility)	1194	– News enabled: True	1233
	1195	• News Configuration:	1234
A.3.2 Middle Environment Configuration	1196	– News impact base scale: 0.4	1235
The Middle configuration uses static data but with full resource capacity and all categories enabled.	1197	– News daily count: 20	1236
	1198	– News random seed: 42	1237
• Time Range: Same as Easy	1199	– News sample ratios:	1238
• Financial Parameters:	1200	* Neutral: 0.9	1239
– Initial funds: 50,000	1201	* Single category: 0.02	1240
– Daily rent: 1,000	1202	* Macro all: 0.03	1241
– Inventory capacity: 40,000	1203	* SKU level: 0.05	1242
• Feature Enablement:	1204	– News impact mode weights:	1243
– Review enabled: True	1205	* Neutral: 0.0	1244
– Review ratio: 0.02	1206	* Macro all: 1.0	1245
– News enabled: False		* Single category: 1.0	1246
• Selected Categories (20 categories):		* SKU level: 1.2	1247
– Bathroom_Tissues, Beer, Bottled_Juices,		• Selected Categories: Same 20 categories as Middle	1248
Canned_Soup, Canned_Tuna	1208		1249
– Cereals, Cheeses, Cigarettes, Cookies,	1209	• Category Effects: All categories have a uniform effect of -0.2	1250
Crackers	1210		1251
– Dish_Detergent, Fabric_Softeners,	1212	• Random Seed: 42	1252
Front_end_candies, Frozen_Entrees	1213		
		A.4 Environment Simulation	1253
		See in Figure ??	1254




Figure 3: Net worth and available funds trajectories of the heuristic policy under different environment configurations, illustrating the calibrated difficulty levels. The *Middle* and *Hard* settings involve a larger number of product categories, enabling higher potential net worth and cash accumulation over the course of an episode.

B Rollout Details

B.1 Hallucinations in Reasoning and Planning

This appendix provides representative examples of hallucinations observed during long-horizon rollouts, focusing on reasoning-time errors that propagate into multi-step planning despite internally coherent logic.

B.1.1 Non-existent SKUs

During execution, models occasionally reference or plan around SKUs that do not exist in the environment. Across all rollouts, we identify 14 distinct non-existent SKUs that repeatedly appear in final daily strategies, indicating persistent hallucinations rather than isolated parsing or formatting errors. Representative examples include SKU identifiers such as 10700013100, 166051312, and 440004627.

Example. In the following strategy output, the model includes a non-existent SKU (440004627) in its set of focus SKUs. This hallucinated identifier is subsequently treated as a valid product in

downstream reasoning and planning:

```
{
  "day": 14,
  "current_date": "1991-09-22",
  "strategy": {
    "macro_strategy": ***,
    "execute_strategy": {
      "focus_skus": [
        "3000001460",
        "3700063037",
        "440004627",
        "5100001251"
      ],
      "news_to_monitor": ***,
    },
    "today_action": ***,
  }
}
```

B.1.2 Hallucinated Dates

Models also hallucinate temporal information during tool-based reasoning when required to provide calendar dates that are not explicitly specified by the environment. Instead of querying the current simulation date, the model resolves this underspecification by fabricating a plausible calendar mapping in order to proceed with execution.

Example. The following excerpt illustrates the model’s reasoning process when attempting to issue a tool call that requires valid date strings:

Tool requires dates in YYYY-MM-DD or MM/DD/YY.

Only day index is known: Day 14.

No calendar start date is provided.

Assume Day 1 = 2023-01-01.

Infer Day 14 = 2023-01-14.

start_date = 2023-01-13

end_date = 2023-01-14

Based on this fabricated assumption, the model issues a syntactically valid tool call using the inferred dates. Although the reasoning chain is internally consistent, the assumed calendar mapping is not grounded in the true environment state, resulting in a semantically incorrect interaction.

These examples highlight a recurring failure mode in which models resolve underspecified variables through plausible fabrication rather than explicit information acquisition, leading to planning

trajectories that appear coherent yet diverge from the actual environment dynamics.

B.2 Invalid or Irrational Actions

We identify invalid or economically irrational actions by applying a set of heuristic rules to model-issued tool calls. Specifically, we flag actions that violate basic economic or operational constraints, including: (i) setting product prices to zero, negative values, or unrealistically high levels (e.g., exceeding 50); and (ii) placing orders with quantities far beyond plausible operational scales for a single SKU.

Across all evaluated rollouts, we detect more than 300 such anomalous tool calls. Among them, 197 correspond to invalid or irrational price modifications, while 125 involve excessively large ordering decisions. These behaviors are observed across multiple models and environment configurations.

Excessive Ordering. The following example is taken from a rollout of the *Kimi K2 Thinking* model in the *Hard* environment. The model places an implausibly large order for a single SKU, far exceeding realistic replenishment volumes:

```
tool: place_order
model: Kimi K2 Thinking
sku_id: 5100000011
quantity: 18000
```

Although syntactically valid, such actions introduce extreme inventory shocks that are misaligned with realistic retail operations.

Invalid Price Modifications. Irrational pricing behavior is also observed across models. In the following example, the model updates the price of a SKU by several orders of magnitude (log excerpted from `still_hard/kimi_thinking/tool_calls.jsonl`):

```
tool: modify_sku_price
model: Kimi K2 Thinking
old_price: 0.25
new_price: 999.0
```

In other cases, models attempt to assign zero or negative prices, which are explicitly rejected by the environment. While such invalid actions are prevented at execution time, extreme but valid values can still propagate downstream and destabilize subsequent decision-making.

Overall, these results indicate that current LLM-based agents lack reliable internal mechanisms for enforcing basic economic plausibility at the action




Figure 4: Macro strategy similarity over time in the easy environment. Higher values indicate greater consistency in high-level decisions across days.



Figure 5: Execution strategy similarity over time in the easy environment. Execution-level behaviors exhibit substantially higher temporal variability than macro strategies.

level, even when tool interfaces enforce syntactic validity and detailed execution logs are available for inspection.

B.3 Rollout Results

See in Table ??

B.4 Strategy Score Analysis

See in Figure ??, ??

B.5 Tool use Analysis

See in Figure ??

B.6 Category Analysis

See in Table ??

B.7 All Model and Framework Results




Figure 6: Correlation between the frequency of SKU review queries during rollouts and average daily sales across all runs. Each point represents a single rollout. We observe a positive association between review-related tool usage and sales performance. The Pearson correlation coefficient r and the corresponding p -value are reported in the figure.

See in Table ??

C Prompt

C.1 Evolving Strategy & Execution Framework Evolving Strategy Phase System Prompts

You are a retail strategy analyst. Your task is to analyze current business data and determine whether the current strategy needs adjustment.

Environment Characteristics

- The store operates with a large number of SKUs, where products within the same category interact and may substitute or cannibalize each other's demand.
 - Historical sales data provides essential signals for future decision-making.
 - Customer reviews dynamically influence product demand and sales velocity, with recent reviews having stronger effects.
- {news_characteristic}
- Supply chains involve delivery lead times, requiring forward-looking inventory planning.

- Orders require delivery time: When you place an order (place_order), the items will not arrive immediately. The delivery time varies and can take up to 7 days (within 7 days). You should plan your inventory accordingly and account for this lead time when making ordering decisions. Orders placed today will arrive within 7 days, but the exact arrival time is variable.
- Inventory items depreciate in value over time and may require disposal when approaching expiration.
- Supplier heterogeneity affects product quality perceptions and customer reviews, leading to supplier-dependent demand outcomes.
- Daily rent: The store incurs a fixed daily rent cost of {daily_rent} that must be paid each day. This daily operating cost is automatically deducted at the end of each day and makes cash-flow management critical for long-term survival and profitability. You must ensure sufficient funds are available to cover the daily rent. The daily rent amount is fixed and must be paid every single day, regardless of sales performance or other factors.

Your Role in Strategy Phase

Each day starts with a STRATEGY PHASE where you:

1. Review the current strategy (provided at the start of this phase)
 2. Use data analysis tools to gather information about:
 - Current inventory status
 - Recent sales history (last 30-60 days)
 - Customer reviews and ratings
 - Supplier prices and quality
- {news_data_point}
- Current financial status
3. Compare current situation with previous days to identify significant changes

4. Set the strategy using the three
 - separate tools (set_macro_strategy,
 - set_execute_strategy, set_action) to
 - set the three strategy components

Strategy Format

The strategy consists of three
→ components:

1. **macro_strategy**: A list of broad
 - strategic guidelines (array of
 - strings)
 - Examples: ["Focus on high-margin
 - products", "Maintain competitive
 - pricing", "Prioritize inventory
 - turnover"]
2. **execute_strategy**: An object with
 - seven fields, all values are arrays:
 - **focus_skus**: Array of SKU IDs
 - that need attention (e.g.,
 - ["SKU_001", "SKU_002"])
 - **sku_supplier_mapping**: Array of
 - mapping objects (e.g.,
 - [{"sku_id": "SKU_001",
 - "supplier_id": "supplier_A"}],
 - [{"sku_id": "SKU_002",
 - "supplier_id": "supplier_B"}])
 - {news_to_monitor_field}
 - **skus_to_reorder**: Array of SKU
 - IDs that need reordering (e.g.,
 - ["SKU_003", "SKU_004"])
 - **price_adjustments**: Array of
 - price adjustment objects (e.g.,
 - [{"sku_id": "SKU_001",
 - "adjustment": "increase by
 - 10%"}], [{"sku_id": "SKU_002",
 - "adjustment": "decrease by
 - 5%"}])
 - **sku_to_monitor**: Array of SKU
 - IDs that should be closely
 - monitored (e.g., ["SKU_005",
 - "SKU_006"])
 - **other**: Array of other strategy
 - notes or metadata (e.g.,
 - [{"comment": "..."}],
 - [{"risk_level": "high"}])

3. **today_action**: An array of action
 - objects, each representing a
 - concrete action using the parameter
 - format of `place\_order` or
 - `modify\_sku\_price`.

- Each action MUST be an object of
 - the form:
 - {"tool": "place_order",
 - "arguments": [{"place_order
 - arguments>}]}
 - OR {"tool": "modify_sku_price",
 - "arguments":
 - [{"modify_sku_price
 - arguments>}]}
 - Example:
 - [
 - { "tool": "place_order",
 - "arguments": { "sku_id":
 - "SKU_001", "supplier_id":
 - "supplier_A", "quantity":
 - 100 } },
 - { "tool": "modify_sku_price",
 - "arguments": { "sku_id":
 - "SKU_002", "new_price":
 - 9.99 } }
 -]

Strategy Setting Tools

Use three separate tools to set

- different parts of the strategy:
- **set_macro_strategy**: Set the
 - macro_strategy (array of strings)
 - Parameter: `macro\_strategy` (array
 - of strings)
 - Example: set_macro_strategy(macro_s_
 - trategy=["Focus on high-margin
 - products", "Maintain competitive
 - pricing"])
- **set_execute_strategy**: Set the
 - execute_strategy (object with seven
 - fields, all arrays)
 - Parameter: `execute\_strategy`
 - (object with fields: focus_skus,
 - sku_supplier_mapping{news_to_mon_
 - itor_param}, skus_to_reorder,
 - price_adjustments, sku_to_monitor,
 - other)
 - All field values must be arrays

- Example: `set_execute_strategy(execute_strategy={{"focus_skus": ["SKU_001"], "sku_supplier_mapping": [{"sku_id": "SKU_001", "supplier_id": "supplier_A"}]}, ...}}`
- ****set_action****: Set the today_action (array of action objects)
 - Parameter: `action` (array of objects, each with "tool" and "arguments" fields)
 - Each action object: `{{"tool": "place_order" | "modify_sku_price", "arguments": {...}}}`
 - Example:
 - `set_action(action=[{"tool": "place_order", "arguments": {"sku_id": "SKU_001", "supplier_id": "supplier_A", "quantity": 100}}])`

You can call these tools multiple times
 → to build or modify the strategy.
 → After your analysis, set all three
 → components to reflect your
 → decisions.

Available Tools for Analysis

The available function signatures are
 → provided within `<tools></tools>` XML
 → tags:
`<tools>`
`{tool_definitions}`
`</tools>`

For each function call, return a JSON
 → object with function name and
 → arguments inside
 → `<tool_call></tool_call>` XML tags:
`<tool_call>`
`{{"name": <function-name>, "arguments":`
 → `<args-json-object>}}`
`</tool_call>`

Important Analysis Tools

Use these tools to gather data:

- `view_funds_and_date`: Check current
 → funds and date
- `view_inventory`: Check current
 → inventory levels
- `view_sku_sales_history`: Analyze sales
 → trends (use last 30-60 days)
- `view_sku_avg_ratings`: Check customer
 → satisfaction
- `view_current_date_supplier_prices`:
 → Check supplier availability and
 → prices
`{news_tools_list}`
- `view_current_orders`: Check pending
 → orders
- `set_macro_strategy`: Set the macro
 → strategy (array of strings)
- `set_execute_strategy`: Set the execute
 → strategy (object with seven fields,
 → all arrays)
- `set_action`: Set the today action
 → (array of action objects)

Note: You CANNOT use `place_order` or
 → `modify_sku_price` in the strategy
 → phase. These tools are only
 → available in the execution phase.

After completing your analysis and
 → updating the strategy, the system
 → will transition to the EXECUTION
 → PHASE.

C.2 Evolving Strategy & Execution Execution Framework Phase System Prompts

1392

1393

You are a retail operations agent
 → executing daily operations based on
 → the current strategy.

Your Role in Execution Phase

In the EXECUTION PHASE, you will receive
 → the ****final strategy**** determined in
 → the Strategy Phase. This strategy
 → includes:

- ****macro_strategy****: Broad strategic
 → guidelines (array of strings)
- ****execute_strategy****: Specific
 → operational details (object with
 → seven fields, all arrays)
- ****today_action****: Concrete actions to
 → take today (array of action objects)

Important Operational Constraints

- ****Daily rent****: The store incurs a
 - ↪ fixed daily rent cost of
 - ↪ {daily_rent} that must be paid each
 - ↪ day. This daily operating cost is
 - ↪ automatically deducted at the end of
 - ↪ each day. Ensure you have sufficient
 - ↪ funds to cover this daily expense.
 - ↪ The daily rent amount is fixed and
 - ↪ must be paid every single day,
 - ↪ regardless of sales performance or
 - ↪ other factors. This makes cash-flow
 - ↪ management critical for long-term
 - ↪ survival and profitability.
- ****Order delivery time****: When you
 - ↪ place an order using place_order,
 - ↪ the items will not arrive
 - ↪ immediately. The delivery time
 - ↪ varies and can take up to 7 days
 - ↪ (within 7 days). Orders placed today
 - ↪ will arrive within 7 days, but the
 - ↪ exact arrival time is variable. Plan
 - ↪ your inventory and ordering
 - ↪ decisions accordingly, considering
 - ↪ the lead time for items to arrive.

Strategy Usage Guidelines

****The strategy is provided as REFERENCE,**
↪ but you can and should make
↪ additional actions based on
↪ real-time data:******

1. ****Reference the strategy**** to
 - ↪ understand priorities and planned
 - ↪ actions:
 - Use macro_strategy for overall
 - ↪ decision-making direction
 - Use execute_strategy fields
 - ↪ (focus_skus, sku_supplier_mappin_
 - ↪ g{news_to_monitor_ref},
 - ↪ skus_to_reorder,
 - ↪ price_adjustments,
 - ↪ sku_to_monitor, other) as
 - ↪ guidance
 - Consider today_action as suggested
 - ↪ actions to take
2. ****Perform additional data queries****
 - ↪ to validate and refine decisions:

- Check current inventory levels,
 - ↪ sales history, supplier
 - ↪ prices{news_impacts_ref}, funds,
 - ↪ etc.
- Use tools like view_inventory,
 - ↪ view_sku_sales_history,
 - ↪ view_current_date_supplier_price_
 - ↪ s{news_tools_ref}, etc.

3. ****Execute actions flexibly****:

- You can execute actions from
 - ↪ today_action when they still make
 - ↪ sense given the latest data
- You can ****adjust, skip, or modify****
 - ↪ actions from today_action if your
 - ↪ analysis shows better
 - ↪ alternatives
- You can ****add additional actions****
 - ↪ beyond today_action if needed
 - ↪ (e.g., unexpected inventory
 - ↪ changes, new supplier
 - ↪ prices{news_impacts_example})
- You can use information from
 - ↪ execute_strategy (like
 - ↪ focus_skus, sku_supplier_mapping)
 - ↪ to make decisions even if not
 - ↪ explicitly in today_action

4. ****End the day**** by calling end_today
 - ↪ when you've completed all operations
 - ↪ for today.

Important Constraints

- You **MUST NOT** modify the stored
 - ↪ strategy itself in this phase
 - ↪ (strategy can only be changed in the
 - ↪ Strategy Phase)
- You **CANNOT** call any tool that changes
 - ↪ macro_strategy / execute_strategy /
 - ↪ today_action
- You **SHOULD** use the strategy as
 - ↪ guidance but make final decisions
 - ↪ based on current data and analysis

Available Tools

The available function signatures are
↪ provided within <tools></tools> XML
↪ tags:
<tools>
{tool_definitions}

</tools>

For each function call, return a JSON
→ object with function name and
→ arguments inside
→ <tool_call></tool_call> XML tags:
<tool_call>
{"name": <function-name>, "arguments":
→ <args-json-object>}}
</tool_call>

Ending the Day

When you have completed all reasonable
→ operations for the day (especially
→ those in today_action, adjusted as
→ needed by current data), you MUST
→ call end_today to advance to the
→ next day. This will trigger a new
→ Strategy Phase for the next day.

C.3 Reflection Framework Reflection Phase Prompts

This prompt is used to generate
→ reflections after each day in
→ \texttt{run_reflection.py}.

\begin{verbatim}
You are a retail operations analyst
→ reflecting on the day's performance.

Task Goal
{task_spec}

Day {day} End Result
{end_today_result.get('formatted',
→ safe_dump(end_today_result))}

Day {day} Interaction History
{interaction_summary}
{memory_context}

Your Task

Generate a comprehensive reflection on
→ today's performance. This reflection
→ should be a complete, detailed
→ analysis that will replace previous
→ reflections. Include:

1. **Performance Summary**: Overall
→ assessment of today's operations,
→ including key metrics (funds,
→ inventory, sales, etc.)
2. **Issue Identification**: What
→ specific problems or challenges
→ occurred? Be specific about what
→ went wrong.
3. **Root Cause Analysis**: Why did
→ these problems happen? Analyze the
→ interaction history to understand
→ what actions or decisions led to the
→ issues. Trace back through the day's
→ operations.
4. **What Worked Well**: Identify any
→ successful strategies or decisions
→ that should be continued.
5. **Actionable Improvements**: What
→ should be done differently next
→ time? Provide specific, actionable
→ recommendations for future
→ operations.
6. **Key Learnings**: What are the most
→ important lessons learned from today
→ that should guide future
→ decision-making?

Format your reflection as a
→ comprehensive, detailed analysis
→ (multiple paragraphs, not just a few
→ sentences). This reflection will be
→ the complete memory used for future
→ days, so it should be thorough and
→ cover all important aspects.

Reflection:

C.4 Macro Similarity Judge Prompt

Please compare the similarity of the
→ following two macro strategies.

Strategy 1's macro_strategy:
{macro1_formatted}

Strategy 2's macro_strategy:
{macro2_formatted}

1396

Please evaluate the similarity between

- ↪ these two strategies and provide a
- ↪ score between 0 and 1, where:
- 1.0 means identical or almost identical
- ↪ identical
- 0.8-0.9 means very similar with only minor differences
- ↪ minor differences
- 0.6-0.7 means somewhat similar with some common points
- ↪ some common points
- 0.4-0.5 means somewhat similar but with significant differences
- ↪ with significant differences
- 0.2-0.3 means not very similar
- 0.0-0.1 means completely different

Please return only a floating-point

- ↪ number between 0 and 1, without any
- ↪ additional text or explanation.

D Agent Framework 1397

D.1 Illustration of Agent Framework 1398

See in Figure ?? 1399




Figure 7: Illustration of Evolving Strategy and Execution Framework

Strategy Type	Strategy Value
macro_strategy	Tissues (SKU 3700060511) face an imminent stockout with current inventory below one day of coverage; incoming replenishments scheduled for 10-08 and 10-10 imply an unavoidable short-term gap and an estimated loss of approximately 1,200 units. Soups (SKU 5100000011) exhibit declining sales and persistently high return rates following a recent price increase, indicating substantial margin erosion. Cigarettes and ginger ale remain operationally stable with high customer ratings and low return rates. Overall operational risk is assessed as high due to anticipated revenue loss and return-driven inefficiencies.
execute_strategy	focus_skus: 3700060511, 5100000011, 1254612128 sku_supplier_mapping: (3700060511, supplier_4), (5100000011, supplier_1), (1230000014, supplier_1), (1690000012, supplier_3), (1254612128, supplier_3) news_to_monitor: [] skus_to_reorder: [] price_adjustments: sku_id = 5100000011, adjustment = decrease to 0.45 or by 20% to stimulate demand sku_to_monitor: 3700060511, 5100000011 other: Tissues — stockout gap 10-05/06/07 estimated 3 days, approximately 1200 units lost at price 0.65 (~780 revenue); incoming supply covers post-10-08; no further reorder feasible Soups — return rate approximately 25% persists despite supplier_1; reviews indicate supplier_4 dominant issues but effects persist; sales declined after price hike; monitor next 3 days sales, returns, and supplier quality Mints — low stock but incoming sufficient Core — stable operations excluding tissues and soups risks; customer traffic approximately 30k/day steady risk_level: high
today_action	place_order: (3700060511, supplier_4, quantity = 800), (1254612128, supplier_3, quantity = 600) modify_sku_price: (sku_id = 5100000011, new_price = 0.45), (sku_id = 1690000012, new_price = 0.78)

Table 6: Example Strategy

Model	Avg. Days $\uparrow$	Avg. Daily Sales $\uparrow$	Avg. Daily Income $\uparrow$	Expiry Ratio $\downarrow$	Return Ratio $\downarrow$	Max Days $\uparrow$
<i>Difficulty: EASY (5 Categories)</i>						
DeepSeek-V3.2 (Exp.)	58.33	229.19	183.26	0.0889	0.1122	66
Gemini-3 (Fast)	50.67	399.39	294.71	0.0799	0.1311	59
GLM-4.6	52.40	174.34	124.67	0.0773	0.1293	58
Grok-4.1 Fast	61.75	508.08	336.94	0.0417	0.0847	88
Kimi-K2 (Thinking)	54.25	260.68	168.72	0.0239	0.1179	58
OpenAI-5.1 Mini	51.75	192.90	122.46	0.0360	0.1237	55
Qwen-235B	37.50	420.31	236.50	0.0745	0.0841	48
Average (7 models)	52.38	301.98	203.30	0.0590	0.1068	61.71
Hand-crafted Policy	180.00	674.18	729.46	0.0266	0.0070	180
<i>Difficulty: MIDDLE (20 Categories)</i>						
DeepSeek-V3.2 (Exp.)	54.67	263.68	255.30	0.1064	0.1818	63
Gemini-3 (Fast)	42.67	439.05	449.04	0.0188	0.1630	48
GLM-4.6	54.33	182.55	131.90	0.0237	0.1274	56
Grok-4.1 Fast	51.00	720.69	398.23	0.0407	0.1160	59
Kimi-K2 (Thinking)	37.00	347.78	356.10	0.0037	0.1677	50
OpenAI-5.1 Mini	56.33	336.24	223.60	0.1307	0.1235	57
Qwen-235B	29.33	216.06	179.42	0.0639	0.1333	45
Average (7 models)	46.48	364.24	282.48	0.0491	0.1399	54.00
Hand-crafted Policy	180.00	1870.21	2809.39	0.0272	0.0074	180
<i>Difficulty: HARD (20 Categories)</i>						
DeepSeek-V3.2 (Exp.)	56.67	165.86	175.79	0.0787	0.1978	59
Gemini-3 (Fast)	35.33	513.10	650.94	0.0906	0.1542	44
GLM-4.6	53.33	205.76	203.07	0.0645	0.1648	55
Grok-4.1 Fast	33.67	504.57	364.52	0.0424	0.1797	50
Kimi-K2 (Thinking)	43.67	248.64	312.82	0.0113	0.1889	57
OpenAI-5.1 Mini	55.33	331.36	335.32	0.1949	0.1923	57
Qwen-235B	40.33	433.64	268.06	0.1746	0.1834	48
Average (7 models)	45.48	320.96	311.28	0.0960	0.1791	52.86
Hand-crafted Policy	180.00	1667.84	2748.94	0.0507	0.0075	180

Table 7: Performance comparison of seven large language models under three difficulty levels. Lower Expiry and Return ratios indicate better operational stability. A hand-crafted policy is included as an approximate upper bound for each difficulty.

Model	Context	Easy		Middle		Hard	
		SKU ↑	Cat ↑	SKU ↑	Cat ↑	SKU ↑	Cat ↑
DeepSeek-V3.2 (Exp.)	128K	7.80	<u>4.07</u>	6.43	4.25	4.71	3.65
Gemini-3 (Fast)	1M	8.50	3.95	11.89	11.18	13.32	8.09
GLM-4.6	200K	6.61	3.55	6.01	3.48	<u>8.47</u>	<u>6.05</u>
OpenAI-5.1 Mini	400K	6.80	2.64	<u>6.82</u>	<u>6.81</u>	7.87	5.19
Grok-4.1 Fast	200K	<u>7.88</u>	4.35	6.51	3.53	7.69	3.90
Kimi-K2 (Thinking)	256K	5.92	2.94	5.53	4.09	6.03	3.22
Qwen-235B (Thinking)	256K	4.78	3.66	6.49	4.05	5.77	4.23
Heuristic Strategy	–	25	5	96	20	96	20

Table 8: SKU and Category counts observed during the strategy stage across difficulties. Context denotes the maximum supported context length of each model. **Bolded** indicates the best model; underlined indicates the second best (Heuristic Strategy excluded).

Model	Avg. Days $\uparrow$	Avg. Daily Sales $\uparrow$	Avg. Daily Income $\uparrow$	Expiry Ratio $\downarrow$	Return Ratio $\downarrow$	Max Days $\uparrow$
<i>Framework: Evolving Strategy & Execution</i>						
DeepSeek-V3.2 (Exp.)	58.33	229.19	183.26	0.0889	0.1122	66
Gemini-3 (Fast)	50.67	399.39	294.71	0.0799	0.1311	59
GLM-4.6	52.40	174.34	124.67	0.0773	0.1293	58
Grok-4.1 Fast	61.75	508.08	336.94	0.0417	0.0847	88
Kimi-K2 (Thinking)	54.25	260.68	168.72	0.0239	0.1179	58
OpenAI-5.1 Mini	51.75	192.90	122.46	0.0360	0.1237	55
Qwen-235B (Thinking)	37.50	420.31	236.50	0.0745	0.0841	48
GPT-5.2	81.00	457.21	358.27	0.0660	0.1141	81
Average (8 models)	55.96	330.26	228.19	0.0610	0.1121	64.13
<i>Framework: Reflection (Day-Level)</i>						
DeepSeek-V3.2 (Exp.)	53.00	235.01	170.04	0.0382	0.1200	66
Gemini-3 (Fast)	45.67	447.74	255.38	0.0682	0.1350	50
GLM-4.6	55.00	160.70	125.67	0.0194	0.1176	62
Grok-4.1 Fast	48.33	297.94	197.54	0.1460	0.0925	54
Kimi-K2 (Thinking)	58.33	216.51	184.01	0.0964	0.1255	71
OpenAI-5.1 Mini	53.33	93.04	92.11	0.1062	0.1331	59
Qwen-235B (Thinking)	30.00	371.90	197.53	0.1919	0.1551	43
GPT-5.2	64.00	283.88	260.88	0.1774	0.0887	64
Average (8 models)	50.96	263.34	185.40	0.1055	0.1209	58.63
<i>Framework: Reflection (Step-Level)</i>						
DeepSeek-V3.2 (Exp.)	–	–	–	–	–	–
Gemini-3 (Fast)	–	–	–	–	–	–
GLM-4.6	51.67	92.35	77.18	0.0148	0.1072	57
Grok-4.1 Fast	–	–	–	–	–	–
Kimi-K2 (Thinking)	51.67	181.19	111.29	0.0353	0.1362	53
OpenAI-5.1 Mini	–	–	–	–	–	–
Qwen-235B (Thinking)	–	–	–	–	–	–
GPT-5.2	56.00	398.71	324.01	0.1536	0.1048	56
Average (8 models)	53.11	224.08	170.83	0.0679	0.1161	55.33
<i>Framework: Plan-and-Act</i>						
DeepSeek-V3.2 (Exp.)	–	–	–	–	–	–
Gemini-3 (Fast)	–	–	–	–	–	–
GLM-4.6	48.33	231.35	113.01	0.0198	0.1096	55
Grok-4.1 Fast	–	–	–	–	–	–
Kimi-K2 (Thinking)	48.67	170.26	105.36	0.0000	0.1056	51
OpenAI-5.1 Mini	–	–	–	–	–	–
Qwen-235B (Thinking)	–	–	–	–	–	–
GPT-5.2	64.00	323.88	193.02	0.0152	0.1096	64
Average (8 models)	53.67	241.83	137.13	0.0117	0.1083	56.67
<i>Heuristic Policy (Upper Bound, Easy)</i>						
Hand-crafted Policy	180.00	674.18	729.46	0.0266	0.007	180

Table 9: Performance comparison of eight large language models under four agent frameworks in the EASY environment. A hand-crafted heuristic policy is included as an approximate upper bound.