

RetailBench: Evaluating Long-Horizon Autonomous Decision-Making and Strategy Stability of LLM Agents in Realistic Retail Environments

Anonymous ACL submission

Abstract

Large Language Model (LLM)-based agents have achieved notable success on short-horizon and highly structured tasks, yet their ability to maintain coherent decision-making over long horizons in realistic and dynamic environments remains an open challenge. We introduce *RetailBench*, a high-fidelity benchmark designed to evaluate long-horizon autonomous decision-making in realistic commercial scenarios, where agents must operate under stochastic demand and evolving external conditions.

We further propose the *Evolving Strategy & Execution* framework, which separates high-level strategic reasoning from low-level action execution, enabling adaptive and interpretable strategy evolution over time. This design is crucial for long-horizon tasks, where non-stationary environments and error accumulation require strategies to be revised at a different temporal scale than action execution. Experiments on eight state-of-the-art LLMs across progressively challenging environments show that our framework improves operational stability and efficiency compared to other baselines. However, performance degrades substantially as task complexity increases, revealing fundamental limitations in current LLMs for long-horizon, multi-factor decision-making.

1 Introduction

Recent large language models (LLMs), particularly when augmented with reasoning and tool-use capabilities, have demonstrated strong performance on a variety of cognitively demanding tasks, including code editing, mathematical problem solving, and complex information retrieval (Jimenez et al., 2024; Phan et al., 2025; Gao et al., 2024). However, accumulating empirical evidence indicates that such capabilities fail to generalize into robust, domain-agnostic autonomy, especially in realistic settings requiring long-horizon planning,

persistent objective alignment, and stable behavioral consistency—core prerequisites for participation in real-world economic systems (Amodei, 2024; Kwa et al., 2025; METR, 2025). Correspondingly, existing agent benchmarks—spanning web interaction (Mialon et al., 2023; Wei et al., 2025; Zhou et al., 2024; Deng et al., 2023; Jimenez et al., 2024; Team, 2025b) primarily focus on short-horizon or highly structured tasks, limiting their ability to evaluate sustained interaction with complex, evolving environments. Recent studies on long-horizon autonomy consistently demonstrate that even state-of-the-art agents struggle to maintain coherent strategies over extended time spans (Nof1.ai, 2025; Andon Labs, 2025; Backlund and Petersson, 2025).

To systematically study this challenge, we introduce *RetailBench*, a new benchmark grounded in real-world commercial data and informed by established economic modeling principles. *RetailBench* centers on a supermarket operation scenario that demands long-horizon decision-making, sustained interaction with a dynamic environment, and the integration of heterogeneous historical information. The benchmark evaluates whether LLM-based agents can autonomously sustain realistic business operations under complex, multi-factor conditions.

We evaluate eight state-of-the-art LLMs in this environment. Our results show that current models struggle to maintain stable decision quality as the decision space expands and often fail to incorporate all relevant information. Moreover, hallucinations and economically irrational behaviors frequently emerge during long-horizon execution, leading to environment collapse and preventing sustained autonomous operation.

Our contributions are summarized as follows:

- We introduce *RetailBench*, a high-fidelity benchmark for evaluating long-horizon autonomous

RetailBench Environment




Figure 1: Overview of the hierarchical supermarket environment, illustrating intra-day agent–environment interactions and end-of-day state transition dynamics.

- decision-making in realistic retail environments.
- We propose the *Evolving Strategy & Execution* agent framework, which improves operational stability compared to a Reflection-based baseline.
- Through extensive experiments, we identify systematic failure modes of current LLM-based agents in long-horizon, multi-factor decision-making settings.

2 Environment Construction

2.1 Problem Formulation and Overview

We model supermarket operations as a Markov Decision Process (MDP), in which an autonomous agent manages a single retail store over a finite horizon of days. At each day t , the agent makes a sequence of operational decisions that jointly determine the store’s daily outcomes.

Formally, the MDP is defined as $(\mathcal{S}, \mathcal{A}, \mathcal{T}, R, \gamma)$, where $\mathcal{S}$ denotes the state space, $\mathcal{A}$ the action space, $\mathcal{T}$ the stochastic transition dynamics, R the reward function, and $\gamma \in (0, 1]$ the discount factor.

At the beginning of each day t , the agent observes the initial state $S_{t,0}$ and executes a sequence of intra-day actions indexed by k . Each action $a_{t,k} \in \mathcal{A}$ induces a transition to the next intra-day state $S_{t,k+1}$. When the agent issues the *end-today*

action, the environment transitions to the initial state of the next day, $S_{t+1,0}$, according to the transition dynamics $\mathcal{T}$. Figure 1 illustrates the overall interaction process.

The environment supports long-horizon operation over more than one thousand simulated days, with episodes terminating if the store fails to pay rent for five consecutive days.

2.2 State Space

The intra-day state $S_{t,k}$ summarizes the complete operational context of the store at step k of day t and is composed of multiple interdependent components: $S_{t,k} = (S_{t,k}^{\text{prod}}, S_{t,k}^{\text{inv}}, S_{t,k}^{\text{sup}}, S_{t,k}^{\text{dem}}, S_{t,k}^{\text{ext}}, S_{t,k}^{\text{fin}})$.

- $S_{t,k}^{\text{prod}}$ and $S_{t,k}^{\text{inv}}$ encode product-level attributes and on-hand inventory status, including prices, shelf life, and historical sales records. Product demand is grounded in real-world retail data derived from the Dominick’s dataset ([Kilts Center for Marketing](#)).
- $S_{t,k}^{\text{sup}}$ represents the supply chain state, including supplier prices, quality levels, and delivery lead times, constructed to reflect empirically observed price–quality relationships ([Grewal et al., 2014](#)).
- $S_{t,k}^{\text{dem}}$ captures demand-side signals such as recent customer traffic and aggregated review statistics,

which influence consumer purchasing behavior (Fedewa et al., 2021).

- $S_{t,k}^{\text{ext}}$ represents external contextual information, including active news events with market-, category-, or product-level impact scope, synthesized to resemble real-world financial news (ashraq, 2025).
- $S_{t,k}^{\text{fin}}$ denotes the financial state of the store, including available cash and estimated net worth, accounting for inventory depreciation over product shelf life.

Together, these components define a richly structured yet partially stochastic state space. Further details of the environment design are deferred to Appendix A.

2.3 Action Space

At each time step t , the agent selects an action $a_t \in \mathcal{A}$, defined as a tuple of decision components: $a_t = (a_t^{\text{price}}, a_t^{\text{repl}}, a_t^{\text{info}}, a_t^{\text{mem}}, a_t^{\text{end}})$. Each component corresponds to a distinct class of operational decisions:

- a_t^{price} denotes pricing decisions for individual SKUs;
- a_t^{repl} denotes inventory replenishment decisions, including supplier selection and order quantities;
- a_t^{info} denotes information acquisition actions, such as querying sales history, customer reviews, supplier conditions, or current news;
- a_t^{mem} denotes memory operations that allow the agent to write and retrieve persistent notes across days;
- a_t^{end} denotes a day-termination action that concludes the current decision phase.

Within a single day, the agent may execute multiple information queries and operational adjustments before issuing the a_t^{end} action to advance the environment to the next day.

2.4 Day Transition Dynamics

After the agent issues the day-termination action, the environment transitions to the next day according to $S_{t,k+1} \sim \mathcal{T}(\cdot \mid S_{t,k}, a_t)$, where the transition function $\mathcal{T}$ factorizes into a sequence of stochastic updates that jointly model daily supermarket operations.

Specifically, each day-level transition consists of the following steps:

1. **Customer traffic sampling.** The total customer traffic for the day is sampled according to stochastic demand dynamics.
2. **Sales realization.** Given the sampled customer traffic and exogenous factors (e.g., promotions, seasonal effects, and news signals), the realized sales volume of each product is determined.
3. **Reviews and returns generation.** Customer reviews and product return events are generated conditional on sales outcomes, product quality, and external influences.
4. **Inventory update.** Sold units are deducted from on-hand inventory, and replenishment orders scheduled to arrive on the current day are added to stock.
5. **Financial and exogenous state update.** The agent’s financial state is updated based on realized revenues and costs, after which new exogenous information (e.g., news or market signals) for the next day is generated.

3 Evolving Strategy & Execution Framework

3.1 Framework Detail

Prior agent frameworks, including ReAct, Reflection, and Plan-and-Act (Erdogan et al., 2025a; Yao et al., 2023b), either do not maintain a persistent global strategy, leading to inconsistent behaviors, or revise strategies during action execution, which can induce oscillation and gradual goal drift in long-horizon environments.

To address these limitations, we propose an explicit two-stage interaction framework, termed *Evolving Strategy & Execution*, that separates strategic deliberation from operational execution. The framework maintains a persistent global strategy and restricts strategy updates to a day-level granularity to prevent excessive short-term fluctuations.

In the *Evolving Strategy* stage, the agent may invoke observation and analysis tools to examine environmental feedback and historical outcomes, but cannot execute actions that directly modify the environment. Based on this information, it determines whether to revise the inherited strategy by

Model	Avg. Days $\uparrow$	Avg. Daily Sales $\uparrow$	Avg. Daily Income $\uparrow$	Expiry Ratio $\downarrow$	Return Ratio $\downarrow$	Max Days $\uparrow$
<i>Framework: Evolving Strategy & Execution</i>						
GLM-4.6	52.40	174.34	124.67	0.0773	0.1293	58
Kimi-K2 (Thinking)	54.25	260.68	168.72	0.0239	0.1179	58
GPT-5.2	81.00	457.21	358.27	0.0660	0.1141	81
Average (3 models)	62.55	297.41	217.22	0.0557	0.1204	65.67
<i>Framework: Reflection (Day-Level)</i>						
GLM-4.6	55.00	160.70	125.67	0.0194	0.1176	62
Kimi-K2 (Thinking)	58.33	216.51	184.01	0.0964	0.1255	71
GPT-5.2	64.00	283.88	260.88	0.1774	0.0887	64
Average (3 models)	59.11	220.36	190.19	0.0977	0.1106	65.67
<i>Framework: Reflection (Step-Level)</i>						
GLM-4.6	51.67	92.35	77.18	0.0148	0.1072	57
Kimi-K2 (Thinking)	51.67	181.19	111.29	0.0353	0.1362	53
GPT-5.2	56.00	398.71	324.01	0.1536	0.1048	56
Average (3 models)	53.11	224.08	170.83	0.0679	0.1161	55.33
<i>Framework: Plan-and-Act</i>						
GLM-4.6	48.33	231.35	113.01	0.0198	0.1096	55
Kimi-K2 (Thinking)	48.67	170.26	105.36	0.0000	0.1056	51
GPT-5.2	64.00	323.88	193.02	0.0152	0.1096	64
Average (3 models)	53.67	241.83	137.13	0.0117	0.1083	56.67
<i>Heuristic Policy (Upper Bound, Easy)</i>						
Hand-crafted Policy	180.00	674.18	729.46	0.0266	0.0070	180

Table 1: Performance comparison of three representative large language models under four agent frameworks in the EASY environment. A hand-crafted heuristic policy is included as an approximate upper bound.

adding, refining, or removing components. This stage ensures that long-term intent and planning logic remain explicit rather than implicitly entangled with immediate action choices.

In the *Execution* stage, the strategy is fixed and treated as immutable. The agent generates concrete actions strictly consistent with the current strategy, without further modification. This controlled separation enables clearer attribution between strategy and behavioral outcomes.

By alternating between these stages, the framework enforces a principled decomposition of decision-making, reduces uncontrolled strategy drift, and promotes behavioral stability and interpretability over extended horizons. An illustration is provided in Appendix D.1.

3.2 Hierarchical Policy Representation

To support structured, interpretable, and temporally extended decision-making under the proposed framework, we represent the agent policy using a hierarchical abstraction that separates strategic intent from executable actions. Each policy consists of three conceptual layers:

1. **Macro Strategy**, which captures high-level

managerial objectives that persist across multiple decision steps;

2. **Execution Strategy**, which encodes structured operational guidance in a machine-readable intermediate representation;
3. **Daily Actions**, which specify concrete executable operations issued to the environment.

Detailed policy configurations and example policies are provided in Appendix A.2.1.

4 Experiment Settings

We conduct experiments under three environment configurations with increasing levels of difficulty. These configurations vary in market complexity, budget constraints, and the presence of exogenous dynamics.

4.1 Environment Configurations

We employ a heuristic policy with full access to the environment’s internal state as a calibration baseline for each environment variant. Environment parameters are tuned such that the heuristic policy remains stable across different difficulty levels while still experiencing meaningful operational pressure

in Appendix A.4. To evaluate models’ information-processing and external perception capabilities, we design three environment configurations:

- *Easy*: A controlled environment without dynamic news events or adaptive supplier price–quality relationships. The market contains five product categories. The agent is initialized with a budget of 10,000 and incurs a fixed daily rent of 250.
- *Middle*: A moderately complex environment that expands the product space to all twenty categories while still excluding dynamic news events and supplier adaptations. The initial budget is increased to 50,000, with a daily rent of 1,000.
- *Hard*: The most challenging and realistic environment, incorporating dynamically generated news events and time-varying supplier price–quality relationships. The market includes all twenty product categories. The agent starts with a budget of 50,000, pays a daily rent of 1,000, and receives twenty news items per day.

Detailed specifications of all environment configurations are provided in Appendix A.3.

4.2 Evaluation Metrics

Metrics. We evaluate store-level operational performance using the following metrics ($\uparrow$ indicates higher is better; $\downarrow$ indicates lower is better):

- *Days* ($\uparrow$): the number of operating days before episode termination;
- *MaxDays* ($\uparrow$): the maximum number of operating days achieved across three rollouts;
- *Avg. Daily Sales* ($\uparrow$): the average number of items sold per day;
- *Avg. Daily Income* ($\uparrow$): the average money earned per day;
- *Expiry Ratio* ($\downarrow$): the fraction of products that expire before being sold;
- *Return Ratio* ($\downarrow$): the fraction of sold products that are returned by customers.

All reported metrics are averaged over three independent rollouts, each subject to a fixed maximum execution horizon.

4.3 Experimental Setup

Agent Frameworks. Preliminary experiments indicate that simple ReAct-style interaction frameworks are unstable in long-horizon settings, frequently exhibiting premature episode termination during mid-horizon execution. This observation motivates the evaluation of agent frameworks that explicitly support sustained strategic control.

We compare our proposed *Evolving Strategy & Execution* framework with step-level Reflection, day-level Reflection, and the Plan-and-Act baseline (Shinn et al., 2023; Erdogan et al., 2025a). For the Reflection framework, we implement two variants with different reflection frequencies: (1) step-level reflection, where the agent reflects after each action, and (2) day-level reflection, where reflection is performed once per day after completing daily actions.

Full prompt specifications for all frameworks are provided in Appendix C.

Models and Evaluation Protocol. We evaluate a diverse set of contemporary large language models, including Qwen-235B (Thinking) (Team, 2025a), Kimi K2 (Thinking) (Team et al., 2025b), GLM-4.6 (Team et al., 2025a), DeepSeek-V3.2-Exp (DeepSeek-AI et al., 2025), Gemini-3-Flash-Preview (Google DeepMind, 2024), Grok-4.1 Fast (xAI, 2025), GPT-5-Mini (OpenAI, 2025a), and GPT-5.2 (OpenAI, 2025b). To manage the substantial token costs of long-horizon rollouts, lower-cost variants are used for closed-source models when available.

Each model is evaluated across three environment configurations, with three independent rollouts per configuration; results are averaged across runs. To isolate the effect of the agent framework, we conduct controlled comparisons between our framework and baseline methods in the *Easy* environment under identical conditions. We further implement a hand-crafted policy based on privileged internal knowledge unavailable to the agents, which serves as an approximate upper bound on performance.

5 Results

5.1 Performance Comparison between Agent Frameworks

Due to computational budget constraints, we conduct a comprehensive four-framework comparison on three representative models: GLM-4.6, Kimi-




Figure 2: Category-level sales and profit per category across Easy, Middle, and Hard environments. Results are shown for three representative models.

K2 (Thinking), and GPT-5.2. As reported in Table 1, *Evolving Strategy & Execution* consistently outperforms alternative agent frameworks on core metrics, achieving higher sales and profit while substantially reducing product expiration rates.

Based on these results, we select the two strongest-performing frameworks for broader evaluation across all eight models. When averaged over all models, our proposed framework demonstrates clear improvements over *Reflection (Day-Level)*.

Despite these gains, a substantial gap remains relative to the hand-crafted heuristic policy, which serves as an approximate upper bound. This gap highlights the current limitations of LLM-based agents in long-horizon decision-making. Notably, Table 1 also suggests that larger closed-source models, such as GPT-5.2, tend to achieve more stable and sustained store operation compared to smaller or open-source counterparts, indicating that model capacity remains an important factor alongside framework design.

Table 8 presents the complete results across all eight evaluated models. Across models, we observe recurring and systematic failure modes in the current environment, suggesting that the performance bottlenecks are structural rather than model-specific.

5.2 Performance across Environments with Varying Difficulty

As environment difficulty increases, all models exhibit consistent performance degradation, including shorter operational durations, higher expiry and return ratios, and reduced long-horizon stability. Although the expansion of product categories in the Middle and Hard settings enables higher aggregate sales and profit for some models, *Sales per Category* and *Profit per Category* decline substan-

Model	Context	Easy		Middle		Hard	
		SKU ↑	Cat ↑	SKU ↑	Cat ↑	SKU ↑	Cat ↑
DeepSeek-V3.2 (Exp.)	128K	4.75	3.31	6.83	5.34	6.64	5.45
Gemini-3 (Fast)	1M	8.72	4.34	13.60	12.43	21.57	13.56
GLM-4.6	200K	4.82	2.98	3.23	2.46	<u>7.08</u>	5.22
OpenAI-5.1 Mini	400K	3.66	1.94	4.10	4.10	6.67	4.79
Grok-4.1 Fast	200K	<u>5.78</u>	<u>3.68</u>	4.87	3.08	5.40	3.20
Kimi-K2 (Thinking)	256K	5.06	3.09	5.17	4.77	6.91	4.75
Qwen-235B (Thinking)	256K	4.59	3.67	3.17	2.36	5.11	4.51
Heuristic Strategy	—	9.03	4.87	35.34	18.61	34.82	18.52

Table 2: SKU and Category counts sold by different models across difficulties. Context denotes the maximum supported context length of each model. **Bolded** indicates the best model; underlined indicates the second best (Human excluded).

tially (Figure 2), indicating persistent challenges in effective resource allocation within increasingly high-dimensional decision spaces.

The relatively small performance gap between the Middle and Hard environments can be partly attributed to the delayed impact of intensified news dynamics, whose effects typically unfold over longer time horizons. Appendix B.3 reports the full results across all environments under the proposed framework. Overall, while models demonstrate limited adaptability as task complexity increases, their performance remains substantially below the heuristic upper bound, highlighting persistent limitations in robust long-horizon decision-making under complex and dynamic conditions.

6 Analysis

Based on both quantitative evaluation results and manual inspection of trajectories, we identify several key factors that explain why current models fail to operate reliably in the environment.

We analyze these failure modes from four complementary perspectives.

6.1 Non-scalable Decision-Making Capability

As shown in Appendix B.3, models achieve reasonable performance in the Easy setting but exhibit consistent degradation as the environment scales in complexity. To better understand this phenomenon, we analyze both the average number of stock keeping units (SKUs) and product categories sold per day, as well as the number of SKUs and categories explicitly considered during the *Evolving Strategy* phase (Tables 2 and Appendix B.6).

Across most models, decision-making capability does not scale proportionally with the size of the environment. Instead, performance remains relatively flat despite a substantial increase in the number of available options. Notably, Gemini-3

Model	Supplier	Inventory	Return	Rating	Price	Review	Sales	History
DeepSeek-V3.2 (Exp.)	41.4	100.0	26.7	75.4	6.6	0.6	89.2	0.0
Gemini-3 (Fast)	69.6	100.0	41.0	82.3	67.7	6.3	87.6	2.9
GLM-4.6	93.9	98.3	79.8	77.8	84.1	5.1	96.3	0.0
OpenAI-5.1 Mini	58.8	99.0	5.1	78.6	3.9	10.6	84.7	0.3
Grok-4.1 Fast	89.9	99.8	84.0	94.0	88.6	36.0	96.4	0.3
Kimi-K2 (Thinking)	57.2	90.4	25.9	51.4	36.2	11.0	64.0	0.5
Qwen-235B (Thinking)	76.8	78.3	31.5	58.2	19.8	23.6	74.6	0.0
Average	69.7	95.1	42.0	74.0	43.8	13.3	84.7	1.0

Table 3: Percentage of days on which each model queries specific information sources when making decisions for SKUs included in the final daily strategy. Higher values indicate greater reliance on the corresponding data source.

(Fast), which supports the largest maximum context window, performs better than other models in this respect, suggesting that larger context capacity helps retain salient information even when the effective interaction context is constrained.

Nevertheless, even the strongest models fail to cover the full decision space. This indicates that current systems are unable to expand their effective decision-making scope as the environment grows, leading to systematic performance degradation in larger and more complex settings.

6.2 Incomplete Decision-Making Due to Limited Information Coverage

We analyze the information sources that models attend to when making operational decisions by examining the data queried for the set of SKUs included in each day’s final strategy. This analysis reveals a clear concentration of attention on a limited subset of signals.

Specifically, most models primarily rely on supplier prices, inventory levels, SKU ratings, and historical sales records when deciding how to operate selected SKUs in Table 3. In contrast, several other critical signals—such as recent customer reviews, return rates, and current selling prices—are consistently underutilized or entirely ignored.

Further correlation analysis shows a strong positive relationship between the frequency of SKU reviews queries and average daily sales performance in Appendix B.5. This observation aligns with the underlying environment dynamics and indicates that incomplete information coverage is a key factor limiting decision quality. Overall, these results suggest that models often fail to perform sufficiently comprehensive information gathering, leading to systematically suboptimal operational decisions.

Model	Macro Strategy			Execution Strategy		
	Std_diff (↓)	MAC (↓)	TV (↓)	Std_diff (↓)	MAC (↓)	TV (↓)
DeepSeek-V3.2(Exp.)	0.130	0.090	5.00	0.289	0.144	7.93
Gemini-3 (Fast)	0.136	0.085	3.82	0.293	0.225	10.18
GLM-4.6	0.131	0.096	4.52	0.264	0.209	9.87
OpenAI-5.1 Mini	0.069	0.051	2.50	0.229	0.166	8.00
Grok-4.1 Fast	0.079	0.045	2.71	0.204	0.151	9.39
Kimi K2(Thinking)	0.179	0.133	6.95	0.335	0.271	14.08
Qwen-235B (Thinking)	0.240	0.189	6.51	0.391	0.306	10.57

Table 4: Temporal instability metrics of macro- and execution-level strategy similarity in the Easy environment. All metrics are lower-is-better (↓). Larger values indicate greater temporal instability. Column-wise maximum values are highlighted in bold.

6.3 Temporal Instability in Execution-Level Decision-Making

Beyond limitations in decision scalability and information coverage, we identify temporal instability in execution-level decision-making as a key contributor to long-horizon failure. Even under relatively stable environmental conditions, agents frequently revise their strategies across consecutive days, leading to inconsistent execution trajectories.

Measuring Strategy Similarity. We quantify temporal instability by measuring the similarity between strategies on adjacent days at both macro and execution levels. Macro strategy similarity is assessed using an LLM-based prompt that evaluates semantic consistency between consecutive high-level plans. Execution strategy similarity is computed via set-based Jaccard similarity over key fields, including `focus_skus`, `sku_supplier_mapping`, `news_to_monitor`, and `sku_to_monitor`, with the final execution similarity obtained by averaging across fields.

Instability Metrics. To characterize temporal fluctuations, we compute three complementary metrics: the standard deviation of first-order differences (*Std_diff*) to capture short-term volatility, the mean absolute change (*MAC*) to estimate typical day-to-day variation, and total variation (*TV*) to reflect cumulative long-term instability.

Across all three environment configurations, we observe consistent temporal instability in both macro- and execution-level strategies. For clarity, we present detailed results from the Easy environment in Table 4 and Appendix B.4. Despite its reduced complexity, the Easy setting already exhibits pronounced temporal fluctuations. In particular, macro strategies remain relatively stable over time, whereas execution strategies show substantially larger variability across most models. This effect is especially pronounced for Qwen-235B

(Thinking), which also performs poorly across all environments.

Overall, these results indicate that long-horizon failures arise not only from suboptimal strategy formulation, but more fundamentally from the inability to maintain temporally consistent execution policies over extended horizons.

6.4 Hallucinations and Invalid Actions

Finally, manual inspection reveals recurrent failure patterns that directly break planning correctness and action validity. We distinguish two closely related but practically different issues:

Hallucinations in reasoning Models occasionally generate reasoning traces that reference non-existent SKUs or fabricate numerical quantities, and subsequently incorporate these hallucinated elements into multi-step plans (examples are provided in Appendix B.1). As a result, decisions become misaligned with the true environment state, even when the overall planning structure appears internally coherent.

Invalid or irrational actions. Models sometimes output actions that violate basic constraints or are inconsistent with historical demand, such as negative order quantities or implausible pricing in Appendix B.2. Even though state-modifying operations are restricted to only a small set of tools, these invalid actions occur with non-negligible frequency and can quickly destabilize the system in long-horizon operation.

In realistic operational settings, both hallucinations and invalid actions would be unacceptable and could directly trigger system collapse.

7 Related Work

Long-horizon planning benchmarks for LLMs. In parallel, a growing body of benchmarks targets long-horizon planning abilities of large language models across structured and interactive environments. PlanBench focuses on classical plan generation, while WebShop, Mind2Web, and ScienceWorld (Deng et al., 2023; Wang et al., 2022; Yao et al., 2023a) evaluate multi-step interaction, error recovery, and adaptive behavior in dynamic settings. More recent benchmarks—such as HeroBench, OdysseyBench, UltraHorizon (Luo et al., 2025; Anokhin et al., 2025; Wang et al., 2025), and explicitly stress extended, interdependent decision sequences that require persistent memory, hierar-

chical reasoning, and long-term strategy maintenance. Collectively, these benchmarks suggest that while short-horizon tasks are increasingly tractable, robust long-horizon planning and execution remain a central challenge for LLM-based agents.

Long-horizon agent frameworks. Recent advances in agent design address these challenges by introducing structured frameworks that decouple high-level planning from low-level execution. Approaches such as Plan-and-Act (Erdogan et al., 2025b) and EAGLET (Si et al., 2025) adopt planner-executor architectures to support hierarchical decision-making, dynamic replanning, and improved execution stability over long horizons. Extensions to multi-agent settings ELHPlan (Ling et al., 2025) and plan-aware context management frameworks PAACE (Yuksel, 2025) further emphasize task decomposition, explicit strategy representation, and memory-aware context control. Together, these works indicate that scalable long-horizon autonomy increasingly relies on structured planning modules and controlled execution mechanisms rather than monolithic end-to-end policies.

8 Conclusion

This paper introduces RetailBench, a benchmark for evaluating long-horizon autonomous decision-making in realistic retail environments. RetailBench models supermarket operations as a stochastic, multi-factor, and temporally extended process, requiring agents to reason jointly about pricing, inventory, information acquisition, and financial sustainability. We further propose the Evolving Strategy & Execution framework, which decouples high-level strategy evolution from low-level execution to better support long-horizon autonomy.

Experiments across eight state-of-the-art large language models show that our framework improves operational stability and economic performance compared to a Reflection-based baseline. However, performance degrades sharply as environment complexity increases, revealing persistent limitations in scalability, information use, execution stability, and action validity. These results suggest that while structured agent frameworks mitigate some challenges, current LLM-based agents remain far from robust strategy-aware autonomy in complex dynamic environments. RetailBench thus provides a principled testbed for advancing research on long-horizon decision-making and agentic reasoning.

9 Limitations

Despite its realism and scale, this work has several limitations. First, RetailBench focuses on a single-store supermarket setting; while expressive, it does not capture multi-store coordination, competitive markets, or strategic interactions among multiple autonomous agents. Second, although the environment incorporates stochastic demand, news dynamics, and supply-chain delays, it remains a simulation grounded in historical data and simplified economic assumptions, which may not fully reflect the complexities of real-world retail systems.

Third, our evaluation is limited to prompting-based LLM agents without parameter updates or long-term learning across episodes; stronger performance may be achievable through reinforcement learning, fine-tuning, or hybrid neuro-symbolic approaches. Finally, while we identify key failure modes such as hallucinations and economically irrational actions, we do not propose explicit algorithmic mechanisms to enforce economic constraints or factual grounding during execution. Addressing these limitations—through richer environments, multi-agent extensions, learning-based adaptation, and constraint-aware action control—remains an important direction for future research.

References

- Dario Amodei. 2024. Machines of loving grace. <https://www.darioamodei.com/essay/machines-of-loving-grace>. Accessed: 2025-12-01.
- Andon Labs. 2025. Vending-bench 2: A benchmark for long-horizon business simulation. <https://andonlabs.com/evals/vending-bench-2>. Accessed: 2025-12-10.
- Petr Anokhin, Roman Khalikov, Stefan Rebrikov, Viktor Volkov, Artyom Sorokin, and Vincent Bissonnette. 2025. Herobench: A benchmark for long-horizon planning and structured reasoning in virtual worlds. *Preprint*, arXiv:2508.12782.
- ashraq. 2025. financial-news-articles. <https://huggingface.co/datasets/ashraq/financial-news-articles>. Accessed: 2025-12-01.
- Axel Backlund and Lukas Petersson. 2025. Vending-bench: A benchmark for long-term coherence of autonomous agents. *Preprint*, arXiv:2502.15840.
- DeepSeek-AI, Aixin Liu, Aoxue Mei, Bangcai Lin, Bing Xue, Bingxuan Wang, Bingzheng Xu, Bochao Wu, Bowei Zhang, Chaofan Lin, Chen Dong, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenhao Xu, Chong Ruan, Damai Dai, Daya Guo, Dejian Yang, and 245 others. 2025. Deepseek-v3.2: Pushing the frontier of open large language models. *Preprint*, arXiv:2512.02556.
- Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Samuel Stevens, Boshi Wang, Huan Sun, and Yu Su. 2023. Mind2web: Towards a generalist agent for the web. *Preprint*, arXiv:2306.06070.
- Lufti Eren Erdogan, Nicholas Lee, Sehoon Kim, Suhong Moon, Hiroki Furuta, Gopala Anumanchipalli, Kurt Keutzer, and Amir Gholami. 2025a. Plan-and-act: Improving planning of agents for long-horizon tasks. *Preprint*, arXiv:2503.09572.
- Lufti Eren Erdogan, Nicholas Lee, Sehoon Kim, Suhong Moon, Hiroki Furuta, Gopala Anumanchipalli, Kurt Keutzer, and Amir Gholami. 2025b. Plan-and-act: Improving planning of agents for long-horizon tasks. *Preprint*, arXiv:2503.09572.
- Dave Fedewa, Chris Holder, Wynn Teichner, and Ben Wiseman. 2021. Five-star growth: Using online ratings to design better products. *McKinsey & Company*. Accessed: 2025-11-30.
- Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma, Liang Chen, Runxin Xu, Zhengyang Tang, Benyou Wang, Daoguang Zan, Shanghaoran Quan, Ge Zhang, Lei Sha, Yichang Zhang, Xuancheng Ren, Tianyu Liu, and Baobao Chang. 2024. Omni-math: A universal olympiad level mathematic benchmark for large language models. *Preprint*, arXiv:2410.07985.
- Google DeepMind. 2024. Gemini 3 flash (preview). <https://deepmind.google/technologies/gemini/>. Large multimodal language model, preview version.
- Dhruv Grewal, Jens Nordfält, Anne Roggeveen, Rainer Olbrich, and Hans Christian Jansen. 2014. Price-quality relationship in pricing strategies for private labels. *Journal of Product and Brand Management*, 23(6):429–438.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. 2024. Swe-bench: Can language models resolve real-world github issues? *Preprint*, arXiv:2310.06770.
- University of Chicago Booth School of Business Kilts Center for Marketing. Dominick’s dataset. <https://www.chicagobooth.edu/research/kilts/research-data/dominicks>. Accessed: 2025-11-01.
- Thomas Kwa, Ben West, Joel Becker, Amy Deng, Katharyn Garcia, Max Hasin, Sami Jawhar, Megan Kinniment, Nate Rush, Sydney Von Arx, Ryan Bloom, Thomas Broadley, Haoxing Du, Brian Goodrich, Nikola Jurkovic, Luke Harold Miles,

713	Seraphina Nix, Tao Lin, Neev Parikh, and 6 others.	Cao, Xiaohan Zhang, Xuancheng Huang, Yao Wei,	767
714	2025. Measuring ai ability to complete long tasks .	and 152 others. 2025a. Glm-4.5: Agentic, reason-	768
715	<i>Preprint</i> , arXiv:2503.14499.	ing, and coding (arc) foundation models . <i>Preprint</i> ,	769
716		arXiv:2508.06471.	770
717	Shaobin Ling, Yun Wang, Chenyou Fan, Tin Lun Lam,	Kimi Team, Yifan Bai, Yiping Bao, Guanduo Chen, Jia-	771
718	and Junjie Hu. 2025. Elhplan: Efficient long-horizon	hao Chen, Ningxin Chen, Ruijue Chen, Yanru Chen,	772
719	task planning for multi-agent collaboration . <i>Preprint</i> ,	Yuan kun Chen, Yutian Chen, Zhuofu Chen, Jialei	773
720		Cui, Hao Ding, Mengnan Dong, Angang Du, Chen-	774
721	Haotian Luo, Huaisong Zhang, Xuelin Zhang, Haoyu	zhuang Du, Dikang Du, Yulun Du, Yu Fan, and 150	775
722	Wang, Zeyu Qin, Wenjie Lu, Guozheng Ma, Haiying	others. 2025b. Kimi k2: Open agentic intelligence .	776
723	He, Yingsha Xie, Qiyang Zhou, Zixuan Hu, Hongze	<i>Preprint</i> , arXiv:2507.20534.	777
724	Mi, Yibo Wang, Naiqiang Tan, Hong Chen, Yi R.		
725	Fung, Chun Yuan, and Li Shen. 2025. Ultrahori-	Qwen Team. 2025a. Qwen3 technical report . <i>Preprint</i> ,	778
726	zon: Benchmarking agent capabilities in ultra long-	arXiv:2505.09388.	779
727	horizon scenarios . <i>Preprint</i> , arXiv:2509.21766.		
728		The Terminal-Bench Team. 2025b. Terminal-bench: A	780
729	Daniel McFadden. 1974. Conditional logit analysis of	benchmark for ai agents in terminal environments .	781
730	qualitative choice behavior. In Paul Zarembka, editor,		
731	<i>Fontiers in Econometrics</i> , pages 105–142. Academic	Mark D. Uncles. 1987. Discrete choice analysis: Theory	782
732	press, New York.	and application to travel demand . <i>Journal of the</i>	783
733		<i>Operational Research Society</i> , 38(4):370–371.	784
734	METR. 2025. Measuring ai ability to complete long		
735	tasks . METR blog.	Ruoyao Wang, Peter Jansen, Marc-Alexandre Côté, and	785
736		Prithviraj Ammanabrolu. 2022. Scienceworld: Is	786
737	Grégoire Mialon, Clémentine Fourrier, Craig Swift,	your agent smarter than a 5th grader? <i>Preprint</i> ,	787
738	Thomas Wolf, Yann LeCun, and Thomas Scialom.	arXiv:2203.07540.	788
739	2023. Gaia: a benchmark for general ai assistants .		
740	<i>Preprint</i> , arXiv:2311.12983.	Weixuan Wang, Dongge Han, Daniel Madrigal Diaz,	789
741		Jin Xu, Victor Rühle, and Saravan Rajmohan.	790
742	Nof1.ai. 2025. Alpha arena — exploring the limits	2025. Odysseybench: Evaluating llm agents on	791
743	of large language models as quant traders. https://nof1.ai/blog/TechPost1 . Accessed: 2025-12-	long-horizon complex office application workflows .	792
744	10.	<i>Preprint</i> , arXiv:2508.09124.	793
745			
746	OpenAI. 2025a. Gpt-5 mini. https://platform.	Jason Wei, Zhiqing Sun, Spencer Papay, Scott McK-	794
747	openai.com/docs/models/gpt-5-mini . Large	inney, Jeffrey Han, Isa Fulford, Hyung Won Chung,	795
748	language model — cost-efficient GPT-5 variant.	Alex Tachard Passos, William Fedus, and Amelia	796
749		Glaese. 2025. Browsecomp: A simple yet chal-	797
750	OpenAI. 2025b. Introducing gpt-5.2 . Accessed: 2026-	lenging benchmark for browsing agents . <i>Preprint</i> ,	798
751	02-19.	arXiv:2504.12516.	799
752			
753	Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li,	xAI. 2025. Grok 4.1 fast and agent tools api. https://	800
754	Josephina Hu, Hugh Zhang, Chen Bo Calvin Zhang,	x.ai/news/grok-4-1-fast . Accessed [Your Ac-	801
755	Mohamed Shaaban, John Ling, Sean Shi, Michael	cess Date].	802
756	Choi, Anish Agrawal, Arnav Chopra, Adam Khoja,		
757	Ryan Kim, Richard Ren, Jason Hausenloy, Oliver	Shunyu Yao, Howard Chen, John Yang, and Karthik	803
758	Zhang, Mantas Mazeika, and 1093 others. 2025. Hu-	Narasimhan. 2023a. Webshop: Towards scalable	804
759	manity’s last exam . <i>Preprint</i> , arXiv:2501.14249.	real-world web interaction with grounded language	805
760		agents . <i>Preprint</i> , arXiv:2207.01206.	806
761	Noah Shinn, Federico Cassano, Edward Berman, Ash-		
762	win Gopinath, Karthik Narasimhan, and Shunyu Yao.	Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak	807
763	2023. Reflexion: Language agents with verbal rein-	Shafraan, Karthik Narasimhan, and Yuan Cao. 2023b.	808
764	forcement learning . <i>Preprint</i> , arXiv:2303.11366.	React: Synergizing reasoning and acting in language	809
765		models . <i>Preprint</i> , arXiv:2210.03629.	810
766			
767	Shuzheng Si, Haozhe Zhao, Kangyang Luo, Gang	Kamer Ali Yuksel. 2025. Paace: A plan-aware	811
768	Chen, Fanchao Qi, Minjia Zhang, Baobao Chang,	automated agent context engineering framework .	812
769	and Maosong Sun. 2025. A goal without a plan	<i>Preprint</i> , arXiv:2512.16970.	813
770	is just a wish: Efficient and effective global plan-		
771	ner training for long-horizon agent tasks . <i>Preprint</i> ,	Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou,	814
772	arXiv:2510.05608.	Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue	815
773		Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Gra-	816
774	GLM Team, Aohan Zeng, Xin Lv, Qinkai Zheng,	ham Neubig. 2024. Webarena: A realistic web envi-	817
775	Zhenyu Hou, Bin Chen, Chengxing Xie, Cunxiang	ronment for building autonomous agents . <i>Preprint</i> ,	818
776	Wang, Da Yin, Hao Zeng, Jiajie Zhang, Kedong	arXiv:2307.13854.	819
777	Wang, Lucen Zhong, Mingdao Liu, Rui Lu, Shulin		

A Environment Configuration Details

A.1 State Space Decomposition

We construct the environment using real-world retail data from the Dominick’s dataset (Kilts Center for Marketing). From the 20 available product categories, we select 96 SKUs with the most complete and informative sales records. The richness of these observations enables reliable modeling of the relationship between pricing decisions and realized demand.

Key Symbols. Let $\mathcal{J}$ denote the set of selected SKUs, with cardinality $|\mathcal{J}| = J$ (where $J = 96$ in our experiments). Each SKU $j \in \mathcal{J}$ belongs to a product category $\text{cat}(j) \in \{1, \dots, 20\}$. For each SKU j , let $\mathcal{K}(j)$ denote its associated supplier set, with $|\mathcal{K}(j)| = 5$.

A.1.1 Product State S_t^{prod}

For each SKU j , we maintain a set of static attributes together with time-varying operational signals:

$$S_{t,j}^{\text{prod}} = (\text{id}_j, \text{desc}_j, L_j, p_{jt}, \mathbf{h}_{j,t}^{\text{sales}}, \mathbf{r}_{j,t}). \quad (1)$$

Here, id_j and desc_j denote the unique identifier and textual description of SKU j , respectively. L_j denotes the shelf life (in days), and p_{jt} is the retail price on day t . The term $\mathbf{h}_{j,t}^{\text{sales}}$ encodes historical sales statistics, while $\mathbf{r}_{j,t}$ summarizes aggregated customer review signals, such as average rating and review volume.

Data grounding. The SKU set and cost priors are constructed from the Dominick’s dataset (Kilts Center for Marketing). Historical sales records from the same source are used to calibrate the demand model.

A.1.2 Inventory State S_t^{inv}

Inventory is represented at the SKU level with age tracking. Let I_{jt} denote the on-hand units of SKU j at the beginning of day t . To support shelf-life constraints and depreciation, we track the arrival time and remaining shelf life of each unit.

Capacity constraint and pending queue. Let Cap denote the total inventory capacity of the store. If incoming replenishment orders would violate this constraint, excess units are placed into a first-in-first-out (FIFO) pending queue $\mathcal{Q}_t$ and become available only after sufficient inventory space is

released:

$$\sum_{j \in \mathcal{J}} \sum_{a=0}^{L_j} I_{jt}^{(a)} \leq \text{Cap}. \quad (2)$$

Stockout signal. When realized demand exceeds the available sellable inventory of SKU j on day t , a stockout indicator is triggered. This signal is recorded at the end of day t and exposed to the agent as explicit operational feedback.

Expiration and destruction. At each day transition, inventory units whose residence time exceeds their shelf life are removed from the system.

A.1.3 Supply Chain State S_t^{sup}

Each SKU j is associated with a set of suppliers $k \in \mathcal{K}(j)$. The supplier state includes procurement price c_{jk} , quality level q_{jk} , and lead-time distribution:

$$S_{t,j}^{\text{sup}} = \{(c_{jk}, q_{jk}, \mathcal{L}_{jk}) : k \in \mathcal{K}(j)\}. \quad (3)$$

Lead time $\ell_{jk} \sim \mathcal{L}_{jk}$ is sampled when an order is placed. Procurement prices are discretized into tiers using Dominick’s cost signals (Kilts Center for Marketing), following empirically observed positive correlations between price tier and quality level (Grewal et al., 2014).

Enforced diversity. To avoid degenerate supplier configurations, we enforce that each SKU has (i) one supplier with maximal quality, (ii) one supplier with minimal price and minimal quality, and remaining suppliers spanning intermediate price–quality levels.

A.1.4 Demand Signals and Demand Generation

Demand-related components capture both observable market signals and the stochastic process governing realized demand. We define

$$S_t^{\text{dem}} = (N_t, \{\mathbf{r}_{j,t}\}_{j \in \mathcal{J}}), \quad (4)$$

where N_t denotes daily customer traffic.

Customer traffic. Daily customer traffic is derived from the Dominick’s dataset.

Consumer choice and demand generation. Given customer traffic N_t , prices $\{p_{jt}\}_{j \in \mathcal{J}}$, review signals $\{\mathbf{r}_{j,t}\}$, and external news $\mathcal{E}_t$, we generate realized demand using a discrete-choice model. Following standard practice in empirical economics,

we adopt a Multinomial Logit (MNL) framework (McFadden, 1974; Uncles, 1987).

For SKU j , the raw utility is defined as

$$U_{jt}^{\text{raw}} = \alpha_j + \beta_j p_{jt} + \varepsilon_{jt}, \quad (5)$$

where α_j captures intrinsic preference, β_j models price sensitivity, and ε_{jt} represents idiosyncratic shocks.

We further incorporate review effects, news signals, and within-category substitution:

$$\begin{aligned} \tilde{U}_{jt} &= U_{jt}^{\text{raw}} + \Delta_j(\mathbf{r}_{j,t}, \mathcal{E}_t) \\ &+ \sum_{\substack{i \neq j \\ \text{cat}(i) = \text{cat}(j)}} \gamma_{ji} \exp(\tilde{U}_{it}). \end{aligned} \quad (6)$$

With the outside option utility normalized to zero, the purchase probability for SKU j on day t is

$$p_{jt}^* = \frac{\exp(\tilde{U}_{jt})}{1 + \sum_{i=1}^J \exp(\tilde{U}_{it})}. \quad (7)$$

Realized demand. Given traffic N_t and purchase probabilities $\{p_{jt}^*\}$, potential demand is sampled as

$$y_{jt} \sim \text{Binomial}(N_t, p_{jt}^*), \quad (8)$$

and realized sales are capped by available on-hand inventory.

A.1.5 External Information S_t^{ext} (News Module)

We maintain a set of active news events $\mathcal{E}_t$. Each event $e \in \mathcal{E}_t$ is represented as

$$e = (\text{type}, \text{scope}, \text{target}, \text{side}, \text{sign}, \eta, \text{text}, \text{ttl}) \quad (9)$$

where $\text{scope} \in \{\text{macro}, \text{category}, \text{product}, \text{neutral}\}$, $\text{side} \in \{\text{demand}, \text{supply}, \text{both}\}$, $\text{sign} \in \{+1, -1\}$, η denotes impact magnitude, and ttl is the time-to-live in days. News texts are synthesized using an LLM to resemble financial news corpora (ashraq, 2025).

Impact application. News impacts are incorporated into (i) demand utilities and/or (ii) supplier prices depending on side and scope. Neutral news has $\eta = 0$ by design.

A.1.6 Financial State S_t^{fin}

Let F_t denote available funds at the start of day t . Net worth is computed as

$$\text{NW}_t = F_t + \sum_{j \in \mathcal{J}} \sum_{a=0}^{L_j} I_{jt}^{(a)} \cdot v_j(a), \quad (10)$$

where $v_j(a)$ denotes the per-unit value at age a . We adopt linear shelf-life depreciation:

$$v_j(a) = c_j \cdot \max\left(0, 1 - \frac{a}{L_j}\right), \quad (11)$$

with c_j denoting a reference procurement cost (e.g., the mean or realized supplier cost).

A.2 Policy Detail

A.2.1 Policy Presentation

To support structured and interpretable decision making, we represent the agent policy using a hierarchical abstraction that separates strategic intent from executable actions. Specifically, each policy is composed of three layers:

1. **Macro Strategy**, which captures high-level managerial principles that persist across days;
2. **Execution Strategy**, which encodes structured operational guidance in a machine-readable form;
3. **Daily Actions**, which enumerate concrete executable operations submitted to the environment.

This design enables the agent to reason at different temporal and semantic granularities, while maintaining a clear separation between planning and execution.

Macro Strategy. The macro strategy consists of a set of natural-language statements that describe high-level objectives (e.g., prioritizing inventory turnover or focusing on high-margin products). These statements are non-executable and serve as persistent guidance for downstream decision making.

Execution Strategy. The execution strategy consists of six key components. `focus_skus` specifies the SKUs that require immediate attention. `sku_supplier_mapping` denotes the corresponding suppliers for each SKU. `news_to_monitor` identifies relevant news signals

to track. `skus_to_reorder` indicates SKUs that require replenishment. `price_adjustment` specifies the SKUs whose prices should be adjusted along with the corresponding adjustment magnitudes. `sku_to_monitor` denotes SKUs under observation. The other field captures additional execution directives.

Daily Actions. Daily actions consist of two operation types: `place_order`, which places purchase orders for selected SKUs, and `modify_sku_price`, which adjusts the prices of specified SKUs.

Strategy–Execution Protocol. Policy usage follows a two-phase protocol. In the *strategy phase*, the agent analyzes the environment and constructs its macro and execution strategies. In the subsequent *execution phase*, the finalized strategy is consumed to generate executable daily actions. The strategy is immutable during execution, enforcing a clear separation between deliberation and action.

A.2.2 Policy Example

See in Table 5.

A.3 Environment Setting

A.3.1 Easy Environment Configuration

The Easy configuration is designed for simpler scenarios with limited resources and a reduced category set.

- **Time Range:**

- Data begin time: 06/06/91
- Data end time: 12/31/95
- Store begin time: 09/07/91
- Store ID: 15

- **Financial Parameters:**

- Initial funds: 10,000
- Daily rent: 250
- Inventory capacity: 10,000

- **Feature Enablement:**

- Review enabled: True
- Review ratio: 0.02
- News enabled: False

- **Selected Categories (5 categories):**

- Bathroom_Tissues
- Canned_Soup
- Cigarettes

- Front_end_candies
- Soft_Drinks

- **Category Effects:** All categories have a uniform effect of -0.2

- **Random Seed:** 42 (for reproducibility)

A.3.2 Middle Environment Configuration

The Middle configuration uses static data but with full resource capacity and all categories enabled.

- **Time Range:** Same as Easy

- **Financial Parameters:**

- Initial funds: 50,000
- Daily rent: 1,000
- Inventory capacity: 40,000

- **Feature Enablement:**

- Review enabled: True
- Review ratio: 0.02
- News enabled: False

- **Selected Categories (20 categories):**

- Bathroom_Tissues, Beer, Bottled_Juices, Canned_Soup, Canned_Tuna
- Cereals, Cheeses, Cigarettes, Cookies, Crackers
- Dish_Detergent, Fabric_Softeners, Front_end_candies, Frozen_Entrees
- Frozen_Juices, Oatmeal, Paper_Towels, Snack_Crackers
- Soft_Drinks, Toothpastes

- **Category Effects:** All categories have a uniform effect of -0.2

- **Random Seed:** 42

A.3.3 Hard Environment Configuration

The Hard configuration uses dynamic data with news events enabled, providing the most complex simulation environment.

- **Time Range:** Same as other configurations

- **Financial Parameters:**

- Initial funds: 50,000
- Daily rent: 1,000
- Inventory capacity: 40,000

- **Feature Enablement:**

- Review enabled: True
- Review ratio: 0.02
- News enabled: True

- **News Configuration:**

- News impact base scale: 0.4
- News daily count: 20
- News random seed: 42
- News sample ratios:
 - * Neutral: 0.9
 - * Single category: 0.02
 - * Macro all: 0.03
 - * SKU level: 0.05
- News impact mode weights:
 - * Neutral: 0.0
 - * Macro all: 1.0
 - * Single category: 1.0
 - * SKU level: 1.2

- **Selected Categories:** Same 20 categories as Middle

- **Category Effects:** All categories have a uniform effect of -0.2

- **Random Seed:** 42

A.4 Environment Simulation

See in Figure 3

B Rollout Details

B.1 Hallucinations in Reasoning and Planning

This appendix provides representative examples of hallucinations observed during long-horizon rollouts, focusing on reasoning-time errors that propagate into multi-step planning despite internally coherent logic.

B.1.1 Non-existent SKUs

During execution, models occasionally reference or plan around SKUs that do not exist in the environment. Across all rollouts, we identify 14 distinct non-existent SKUs that repeatedly appear in final daily strategies, indicating persistent hallucinations rather than isolated parsing or formatting errors. Representative examples include SKU identifiers such as 10700013100, 166051312, and 440004627.




Figure 3: Net worth and available funds trajectories of the heuristic policy under different environment configurations, illustrating the calibrated difficulty levels. The *Middle* and *Hard* settings involve a larger number of product categories, enabling higher potential net worth and cash accumulation over the course of an episode.

Example. In the following strategy output, the model includes a non-existent SKU (440004627) in its set of focus SKUs. This hallucinated identifier is subsequently treated as a valid product in downstream reasoning and planning:

```
{
  "day": 14,
  "current_date": "1991-09-22",
  "strategy": {
    "macro_strategy": ***,
    "execute_strategy": {
      "focus_skus": [
        "3000001460",
        "3700063037",
        "440004627",
        "5100001251"
      ],
      "news_to_monitor": ***,
    },
    "today_action": ***,
  }
}
```

B.1.2 Hallucinated Dates

Models also hallucinate temporal information during tool-based reasoning when required to provide calendar dates that are not explicitly specified by the environment. Instead of querying the current simulation date, the model resolves this underspecification by fabricating a plausible calendar mapping in order to proceed with execution.

Example. The following excerpt illustrates the model’s reasoning process when attempting to issue a tool call that requires valid date strings:

```
Tool requires dates in YYYY-MM-DD
or MM/DD/YY.
Only day index is known: Day 14.
No calendar start date is provided.
```

```
Assume Day 1 = 2023-01-01.
Infer Day 14 = 2023-01-14.
```

```
start_date = 2023-01-13
end_date    = 2023-01-14
```

Based on this fabricated assumption, the model issues a syntactically valid tool call using the inferred dates. Although the reasoning chain is internally consistent, the assumed calendar mapping is not grounded in the true environment state, resulting in a semantically incorrect interaction.

These examples highlight a recurring failure mode in which models resolve underspecified variables through plausible fabrication rather than explicit information acquisition, leading to planning trajectories that appear coherent yet diverge from the actual environment dynamics.

B.2 Invalid or Irrational Actions

We identify invalid or economically irrational actions by applying a set of heuristic rules to model-issued tool calls. Specifically, we flag actions that violate basic economic or operational constraints, including: (i) setting product prices to zero, negative values, or unrealistically high levels (e.g., exceeding 50); and (ii) placing orders with quantities far beyond plausible operational scales for a single SKU.

Across all evaluated rollouts, we detect more than 300 such anomalous tool calls. Among them, 197 correspond to invalid or irrational price modifications, while 125 involve excessively large ordering decisions. These behaviors are observed across multiple models and environment configurations.

Excessive Ordering. The following example is taken from a rollout of the *Kimi K2 Thinking* model in the *Hard* environment. The model places an implausibly large order for a single SKU, far exceeding realistic replenishment volumes:

```
tool: place_order
model: Kimi K2 Thinking
sku_id: 5100000011
quantity: 18000
```

Although syntactically valid, such actions introduce extreme inventory shocks that are misaligned with realistic retail operations.

Invalid Price Modifications. Irrational pricing behavior is also observed across models. In the following example, the model updates the price of a SKU by several orders of magnitude (log excerpted from `still_hard/kimi_thinking/tool_calls.jsonl`):

```
tool: modify_sku_price
model: Kimi K2 Thinking
old_price: 0.25
new_price: 999.0
```

In other cases, models attempt to assign zero or negative prices, which are explicitly rejected by the environment. While such invalid actions are prevented at execution time, extreme but valid values can still propagate downstream and destabilize subsequent decision-making.

Overall, these results indicate that current LLM-based agents lack reliable internal mechanisms for enforcing basic economic plausibility at the action level, even when tool interfaces enforce syntactic validity and detailed execution logs are available for inspection.

B.3 Rollout Results

See in Table ??

B.4 Strategy Score Analysis

See in Figure 4, 5

B.5 Tool use Analysis

See in Figure 6

B.6 Category Analysis

See in Table 7

B.7 All Model and Framework Results

s
See in Table 8




Figure 4: Macro strategy similarity over time in the easy environment. Higher values indicate greater consistency in high-level decisions across days.



Figure 5: Execution strategy similarity over time in the easy environment. Execution-level behaviors exhibit substantially higher temporal variability than macro strategies.

C Prompt

C.1 Evolving Strategy & Execution Framework Evolving Strategy Phase System Prompts

You are a retail strategy analyst. Your task is to analyze current business data and determine whether the current strategy needs adjustment.

Environment Characteristics

- The store operates with a large number of SKUs, where products within the same category interact and may substitute or cannibalize each other's demand.




Figure 6: Correlation between the frequency of SKU review queries during rollouts and average daily sales across all runs. Each point represents a single rollout. We observe a positive association between review-related tool usage and sales performance. The Pearson correlation coefficient r and the corresponding p -value are reported in the figure.

- Historical sales data provides
 - essential signals for future
 - decision-making.
- Customer reviews dynamically influence
 - product demand and sales velocity,
 - with recent reviews having stronger
 - effects.
- {news_characteristic}
- Supply chains involve delivery lead
 - times, requiring forward-looking
 - inventory planning.
- Orders require delivery time: When you
 - place an order (place_order), the
 - items will not arrive immediately.
 - The delivery time varies and can
 - take up to 7 days (within 7 days).
 - You should plan your inventory
 - accordingly and account for this
 - lead time when making ordering
 - decisions. Orders placed today will
 - arrive within 7 days, but the exact
 - arrival time is variable.
- Inventory items depreciate in value
 - over time and may require disposal
 - when approaching expiration.
- Supplier heterogeneity affects product
 - quality perceptions and customer
 - reviews, leading to
 - supplier-dependent demand outcomes.

- Daily rent: The store incurs a fixed
 - ↪ daily rent cost of {daily_rent} that
 - ↪ must be paid each day. This daily
 - ↪ operating cost is automatically
 - ↪ deducted at the end of each day and
 - ↪ makes cash-flow management critical
 - ↪ for long-term survival and
 - ↪ profitability. You must ensure
 - ↪ sufficient funds are available to
 - ↪ cover the daily rent. The daily rent
 - ↪ amount is fixed and must be paid
 - ↪ every single day, regardless of
 - ↪ sales performance or other factors.

Your Role in Strategy Phase

Each day starts with a STRATEGY PHASE

↪ where you:

1. Review the current strategy (provided
 - ↪ at the start of this phase)
2. Use data analysis tools to gather
 - ↪ information about:
 - Current inventory status
 - Recent sales history (last 30-60
 - ↪ days)
 - Customer reviews and ratings
 - Supplier prices and quality

{news_data_point}

- Current financial status
3. Compare current situation with
 - ↪ previous days to identify
 - ↪ significant changes
 4. Set the strategy using the three
 - ↪ separate tools (set_macro_strategy,
 - ↪ set_execute_strategy, set_action) to
 - ↪ set the three strategy components

Strategy Format

The strategy consists of three

↪ components:

1. ****macro_strategy****: A list of broad
 - ↪ strategic guidelines (array of
 - ↪ strings)
 - Examples: ["Focus on high-margin
 - ↪ products", "Maintain competitive
 - ↪ pricing", "Prioritize inventory
 - ↪ turnover"]
2. ****execute_strategy****: An object with
 - ↪ seven fields, all values are arrays:

- ****focus_skus****: Array of SKU IDs
 - ↪ that need attention (e.g.,
 - ↪ ["SKU_001", "SKU_002"])
 - ****sku_supplier_mapping****: Array of
 - ↪ mapping objects (e.g.,
 - ↪ [{"sku_id": "SKU_001",
 - ↪ "supplier_id": "supplier_A"}],
 - ↪ [{"sku_id": "SKU_002",
 - ↪ "supplier_id": "supplier_B"}])
- {news_to_monitor_field}

- ****skus_to_reorder****: Array of SKU
 - ↪ IDs that need reordering (e.g.,
 - ↪ ["SKU_003", "SKU_004"])
- ****price_adjustments****: Array of
 - ↪ price adjustment objects (e.g.,
 - ↪ [{"sku_id": "SKU_001",
 - ↪ "adjustment": "increase by
 - ↪ 10%"}], [{"sku_id": "SKU_002",
 - ↪ "adjustment": "decrease by
 - ↪ 5%"}])
- ****sku_to_monitor****: Array of SKU
 - ↪ IDs that should be closely
 - ↪ monitored (e.g., ["SKU_005",
 - ↪ "SKU_006"])
- ****other****: Array of other strategy
 - ↪ notes or metadata (e.g.,
 - ↪ [{"comment": "..."}],
 - ↪ [{"risk_level": "high"}])

3. ****today_action****: An array of action
 - ↪ objects, each representing a
 - ↪ concrete action using the parameter
 - ↪ format of `place\_order` or
 - ↪ `modify\_sku\_price`.

- Each action MUST be an object of
 - ↪ the form:
 - {"tool": "place_order",
 - ↪ "arguments": {{"<place_order
 - ↪ arguments>}}}
 - OR {"tool": "modify_sku_price",
 - ↪ "arguments":
 - ↪ {{"<modify_sku_price
 - ↪ arguments>}}}

- Example:

```
[
  {"tool": "place_order",
   ↪ "arguments": {"sku_id":
   ↪ "SKU_001", "supplier_id":
   ↪ "supplier_A", "quantity":
   ↪ 100}}},
```

```

    {"tool": "modify_sku_price",
    ↪  "arguments": {"sku_id":
    ↪  "SKU_002", "new_price":
    ↪  9.99}}}]
]

```

Strategy Setting Tools

Use three separate tools to set

- ↪ different parts of the strategy:
- ****set_macro_strategy****: Set the macro_strategy (array of strings)
 - ↪ Parameter: `macro\_strategy` (array of strings)
 - ↪ Example: `set_macro_strategy(macro_strategy=["Focus on high-margin products", "Maintain competitive pricing"])`
- ****set_execute_strategy****: Set the execute_strategy (object with seven fields, all arrays)
 - ↪ Parameter: `execute\_strategy` (object with fields: focus_skus, sku_supplier_mapping{news_to_monitor_param}, skus_to_reorder, price_adjustments, sku_to_monitor, other)
 - ↪ All field values must be arrays
 - ↪ Example: `set_execute_strategy(execute_strategy={"focus_skus": ["SKU_001"], "sku_supplier_mapping": [{"sku_id": "SKU_001", "supplier_id": "supplier_A"}], ...})`
- ****set_action****: Set the today_action (array of action objects)
 - ↪ Parameter: `action` (array of objects, each with "tool" and "arguments" fields)
 - ↪ Each action object: `{"tool": "place_order" | "modify_sku_price", "arguments": {...}}`
 - ↪ Example: `set_action(action=[{"tool": "place_order", "arguments": {"sku_id": "SKU_001", "supplier_id": "supplier_A", "quantity": 100}}])`

You can call these tools multiple times

- ↪ to build or modify the strategy.
- ↪ After your analysis, set all three components to reflect your decisions.

Available Tools for Analysis

The available function signatures are

- ↪ provided within `<tools></tools>` XML tags:

```

<tools>
{tool_definitions}
</tools>

```

For each function call, return a JSON

- ↪ object with function name and arguments inside
 - ↪ `<tool_call></tool_call>` XML tags:
- ```

<tool_call>
{"name": <function-name>, "arguments":
↪ <args-json-object>}}
</tool_call>

```

## # Important Analysis Tools

Use these tools to gather data:

- view\_funds\_and\_date: Check current funds and date
  - view\_inventory: Check current inventory levels
  - view\_sku\_sales\_history: Analyze sales trends (use last 30-60 days)
  - view\_sku\_avg\_ratings: Check customer satisfaction
  - view\_current\_date\_supplier\_prices: Check supplier availability and prices
- ```

{news_tools_list}

```
- view_current_orders: Check pending orders
 - set_macro_strategy: Set the macro strategy (array of strings)
 - set_execute_strategy: Set the execute strategy (object with seven fields, all arrays)
 - set_action: Set the today action (array of action objects)

Note: You CANNOT use place_order or
→ modify_sku_price in the strategy
→ phase. These tools are only
→ available in the execution phase.

After completing your analysis and
→ updating the strategy, the system
→ will transition to the EXECUTION
→ PHASE.

- **Order delivery time**: When you
→ place an order using place_order,
→ the items will not arrive
→ immediately. The delivery time
→ varies and can take up to 7 days
→ (within 7 days). Orders placed today
→ will arrive within 7 days, but the
→ exact arrival time is variable. Plan
→ your inventory and ordering
→ decisions accordingly, considering
→ the lead time for items to arrive.

Strategy Usage Guidelines

The strategy is provided as REFERENCE,
→ but you can and should make
→ additional actions based on
→ real-time data:

C.2 Evolving Strategy & Execution Framework Phase System Prompts

You are a retail operations agent
→ executing daily operations based on
→ the current strategy.

Your Role in Execution Phase

In the EXECUTION PHASE, you will receive
→ the **final strategy** determined in
→ the Strategy Phase. This strategy
→ includes:
- **macro_strategy**: Broad strategic
→ guidelines (array of strings)
- **execute_strategy**: Specific
→ operational details (object with
→ seven fields, all arrays)
- **today_action**: Concrete actions to
→ take today (array of action objects)

Important Operational Constraints

- **Daily rent**: The store incurs a
→ fixed daily rent cost of
→ {daily_rent} that must be paid each
→ day. This daily operating cost is
→ automatically deducted at the end of
→ each day. Ensure you have sufficient
→ funds to cover this daily expense.
→ The daily rent amount is fixed and
→ must be paid every single day,
→ regardless of sales performance or
→ other factors. This makes cash-flow
→ management critical for long-term
→ survival and profitability.

1. **Reference the strategy** to
→ understand priorities and planned
→ actions:
 - Use macro_strategy for overall
→ decision-making direction
 - Use execute_strategy fields
→ (focus_skus, sku_supplier_mappin,
→ g{news_to_monitor_ref},
→ skus_to_reorder,
→ price_adjustments,
→ sku_to_monitor, other) as
→ guidance
 - Consider today_action as suggested
→ actions to take
2. **Perform additional data queries**
→ to validate and refine decisions:
 - Check current inventory levels,
→ sales history, supplier
→ prices{news_impacts_ref}, funds,
→ etc.
 - Use tools like view_inventory,
→ view_sku_sales_history,
→ view_current_date_supplier_price,
→ s{news_tools_ref},
→ etc.
3. **Execute actions flexibly**:
 - You can execute actions from
→ today_action when they still make
→ sense given the latest data

- You can ****adjust, skip, or modify****
 - ↪ actions from today_action if your
 - ↪ analysis shows better
 - ↪ alternatives
- You can ****add additional actions****
 - ↪ beyond today_action if needed
 - ↪ (e.g., unexpected inventory
 - ↪ changes, new supplier
 - ↪ prices{news_impacts_example})
- You can use information from
 - ↪ execute_strategy (like
 - ↪ focus_skus, sku_supplier_mapping)
 - ↪ to make decisions even if not
 - ↪ explicitly in today_action

4. ****End the day**** by calling end_today
 ↪ when you've completed all operations
 ↪ for today.

Important Constraints

- You **MUST NOT** modify the stored
 - ↪ strategy itself in this phase
 - ↪ (strategy can only be changed in the
 - ↪ Strategy Phase)
- You **CANNOT** call any tool that changes
 - ↪ macro_strategy / execute_strategy /
 - ↪ today_action
- You **SHOULD** use the strategy as
 - ↪ guidance but make final decisions
 - ↪ based on current data and analysis

Available Tools

The available function signatures are
 ↪ provided within <tools></tools> XML
 ↪ tags:

```
<tools>
{tool_definitions}
</tools>
```

For each function call, return a JSON
 ↪ object with function name and
 ↪ arguments inside
 ↪ <tool_call></tool_call> XML tags:
 <tool_call>
 {{"name": <function-name>, "arguments":
 ↪ <args-json-object>}}
 </tool_call>

Ending the Day

When you have completed all reasonable
 ↪ operations for the day (especially
 ↪ those in today_action, adjusted as
 ↪ needed by current data), you **MUST**
 ↪ call end_today to advance to the
 ↪ next day. This will trigger a new
 ↪ Strategy Phase for the next day.

C.3 Reflection Framework Reflection Phase Prompts

1229

1230

This prompt is used to generate
 ↪ reflections after each day in
 ↪ \texttt{run_reflection.py}.

```
\begin{verbatim}
You are a retail operations analyst
↪ reflecting on the day's performance.
```

```
# Task Goal
{task_spec}
```

```
# Day {day} End Result
{end_today_result.get('formatted',
↪ safe_dump(end_today_result))}
```

```
# Day {day} Interaction History
{interaction_summary}
{memory_context}
```

Your Task

Generate a comprehensive reflection on
 ↪ today's performance. This reflection
 ↪ should be a complete, detailed
 ↪ analysis that will replace previous
 ↪ reflections. Include:

1. ****Performance Summary****: Overall
 - ↪ assessment of today's operations,
 - ↪ including key metrics (funds,
 - ↪ inventory, sales, etc.)
2. ****Issue Identification****: What
 - ↪ specific problems or challenges
 - ↪ occurred? Be specific about what
 - ↪ went wrong.

3. ****Root Cause Analysis****: Why did
 - ↪ these problems happen? Analyze the
 - ↪ interaction history to understand
 - ↪ what actions or decisions led to the
 - ↪ issues. Trace back through the day's
 - ↪ operations.
4. ****What Worked Well****: Identify any
 - ↪ successful strategies or decisions
 - ↪ that should be continued.
5. ****Actionable Improvements****: What
 - ↪ should be done differently next
 - ↪ time? Provide specific, actionable
 - ↪ recommendations for future
 - ↪ operations.
6. ****Key Learnings****: What are the most
 - ↪ important lessons learned from today
 - ↪ that should guide future
 - ↪ decision-making?

Format your reflection as a

- ↪ comprehensive, detailed analysis
- ↪ (multiple paragraphs, not just a few
- ↪ sentences). This reflection will be
- ↪ the complete memory used for future
- ↪ days, so it should be thorough and
- ↪ cover all important aspects.

Reflection:

- 0.4-0.5 means somewhat similar but
- ↪ with significant differences
- 0.2-0.3 means not very similar
- 0.0-0.1 means completely different

Please return only a floating-point

- ↪ number between 0 and 1, without any
- ↪ additional text or explanation.

D Agent Framework

1232

D.1 Illustration of Agent Framework

1233

See in Figure 7

1234

C.4 Macro Similiarity Judge Prompt

Please compare the similarity of the

- ↪ following two macro strategies.

Strategy 1's macro_strategy:
{macro1_formatted}

Strategy 2's macro_strategy:
{macro2_formatted}

Please evaluate the similarity between

- ↪ these two strategies and provide a
- ↪ score between 0 and 1, where:
- 1.0 means identical or almost
- ↪ identical
- 0.8-0.9 means very similar with only
- ↪ minor differences
- 0.6-0.7 means somewhat similar with
- ↪ some common points




Figure 7: Illustration of Evolving Strategy and Execution Framework

Strategy Type	Strategy Value
macro_strategy	Tissues (SKU 3700060511) face an imminent stockout with current inventory below one day of coverage; incoming replenishments scheduled for 10-08 and 10-10 imply an unavoidable short-term gap and an estimated loss of approximately 1,200 units. Soups (SKU 5100000011) exhibit declining sales and persistently high return rates following a recent price increase, indicating substantial margin erosion. Cigarettes and ginger ale remain operationally stable with high customer ratings and low return rates. Overall operational risk is assessed as high due to anticipated revenue loss and return-driven inefficiencies.
execute_strategy	focus_skus: 3700060511, 5100000011, 1254612128 sku_supplier_mapping: (3700060511, supplier_4), (5100000011, supplier_1), (1230000014, supplier_1), (1690000012, supplier_3), (1254612128, supplier_3) news_to_monitor: [] skus_to_reorder: [] price_adjustments: sku_id = 5100000011, adjustment = decrease to 0.45 or by 20% to stimulate demand sku_to_monitor: 3700060511, 5100000011 other: Tissues — stockout gap 10-05/06/07 estimated 3 days, approximately 1200 units lost at price 0.65 (~780 revenue); incoming supply covers post-10-08; no further reorder feasible Soups — return rate approximately 25% persists despite supplier_1; reviews indicate supplier_4 dominant issues but effects persist; sales declined after price hike; monitor next 3 days sales, returns, and supplier quality Mints — low stock but incoming sufficient Core — stable operations excluding tissues and soups risks; customer traffic approximately 30k/day steady risk_level: high
today_action	place_order: (3700060511, supplier_4, quantity = 800), (1254612128, supplier_3, quantity = 600) modify_sku_price: (sku_id = 5100000011, new_price = 0.45), (sku_id = 1690000012, new_price = 0.78)

Table 5: Example Strategy

Model	Avg. Days $\uparrow$	Avg. Daily Sales $\uparrow$	Avg. Daily Income $\uparrow$	Expiry Ratio $\downarrow$	Return Ratio $\downarrow$	Max Days $\uparrow$
<i>Difficulty: EASY (5 Categories)</i>						
DeepSeek-V3.2 (Exp.)	58.33	229.19	183.26	0.0889	0.1122	66
Gemini-3 (Fast)	50.67	399.39	294.71	0.0799	0.1311	59
GLM-4.6	52.40	174.34	124.67	0.0773	0.1293	58
Grok-4.1 Fast	61.75	508.08	336.94	0.0417	0.0847	88
Kimi-K2 (Thinking)	54.25	260.68	168.72	0.0239	0.1179	58
OpenAI-5.1 Mini	51.75	192.90	122.46	0.0360	0.1237	55
Qwen-235B	37.50	420.31	236.50	0.0745	0.0841	48
Average (7 models)	52.38	301.98	203.30	0.0590	0.1068	61.71
Hand-crafted Policy	180.00	674.18	729.46	0.0266	0.0070	180
<i>Difficulty: MIDDLE (20 Categories)</i>						
DeepSeek-V3.2 (Exp.)	54.67	263.68	255.30	0.1064	0.1818	63
Gemini-3 (Fast)	42.67	439.05	449.04	0.0188	0.1630	48
GLM-4.6	54.33	182.55	131.90	0.0237	0.1274	56
Grok-4.1 Fast	51.00	720.69	398.23	0.0407	0.1160	59
Kimi-K2 (Thinking)	37.00	347.78	356.10	0.0037	0.1677	50
OpenAI-5.1 Mini	56.33	336.24	223.60	0.1307	0.1235	57
Qwen-235B	29.33	216.06	179.42	0.0639	0.1333	45
Average (7 models)	46.48	364.24	282.48	0.0491	0.1399	54.00
Hand-crafted Policy	180.00	1870.21	2809.39	0.0272	0.0074	180
<i>Difficulty: HARD (20 Categories)</i>						
DeepSeek-V3.2 (Exp.)	56.67	165.86	175.79	0.0787	0.1978	59
Gemini-3 (Fast)	35.33	513.10	650.94	0.0906	0.1542	44
GLM-4.6	53.33	205.76	203.07	0.0645	0.1648	55
Grok-4.1 Fast	33.67	504.57	364.52	0.0424	0.1797	50
Kimi-K2 (Thinking)	43.67	248.64	312.82	0.0113	0.1889	57
OpenAI-5.1 Mini	55.33	331.36	335.32	0.1949	0.1923	57
Qwen-235B	40.33	433.64	268.06	0.1746	0.1834	48
Average (7 models)	45.48	320.96	311.28	0.0960	0.1791	52.86
Hand-crafted Policy	180.00	1667.84	2748.94	0.0507	0.0075	180

Table 6: Performance comparison of seven large language models under three difficulty levels. Lower Expiry and Return ratios indicate better operational stability. A hand-crafted policy is included as an approximate upper bound for each difficulty.

Model	Context	Easy		Middle		Hard	
		SKU ↑	Cat ↑	SKU ↑	Cat ↑	SKU ↑	Cat ↑
DeepSeek-V3.2 (Exp.)	128K	7.80	<u>4.07</u>	6.43	4.25	4.71	3.65
Gemini-3 (Fast)	1M	8.50	3.95	11.89	11.18	13.32	8.09
GLM-4.6	200K	6.61	3.55	6.01	3.48	<u>8.47</u>	<u>6.05</u>
OpenAI-5.1 Mini	400K	6.80	2.64	<u>6.82</u>	<u>6.81</u>	7.87	5.19
Grok-4.1 Fast	200K	<u>7.88</u>	4.35	6.51	3.53	7.69	3.90
Kimi-K2 (Thinking)	256K	5.92	2.94	5.53	4.09	6.03	3.22
Qwen-235B (Thinking)	256K	4.78	3.66	6.49	4.05	5.77	4.23
Heuristic Strategy	–	25	5	96	20	96	20

Table 7: SKU and Category counts observed during the strategy stage across difficulties. Context denotes the maximum supported context length of each model. **Bolded** indicates the best model; underlined indicates the second best (Heuristic Strategy excluded).

Model	Avg. Days $\uparrow$	Avg. Daily Sales $\uparrow$	Avg. Daily Income $\uparrow$	Expiry Ratio $\downarrow$	Return Ratio $\downarrow$	Max Days $\uparrow$
<i>Framework: Evolving Strategy & Execution</i>						
DeepSeek-V3.2 (Exp.)	58.33	229.19	183.26	0.0889	0.1122	66
Gemini-3 (Fast)	50.67	399.39	294.71	0.0799	0.1311	59
GLM-4.6	52.40	174.34	124.67	0.0773	0.1293	58
Grok-4.1 Fast	61.75	508.08	336.94	0.0417	0.0847	88
Kimi-K2 (Thinking)	54.25	260.68	168.72	0.0239	0.1179	58
OpenAI-5.1 Mini	51.75	192.90	122.46	0.0360	0.1237	55
Qwen-235B (Thinking)	37.50	420.31	236.50	0.0745	0.0841	48
GPT-5.2	81.00	457.21	358.27	0.0660	0.1141	81
Average (8 models)	55.96	330.26	228.19	0.0610	0.1121	64.13
<i>Framework: Reflection (Day-Level)</i>						
DeepSeek-V3.2 (Exp.)	53.00	235.01	170.04	0.0382	0.1200	66
Gemini-3 (Fast)	45.67	447.74	255.38	0.0682	0.1350	50
GLM-4.6	55.00	160.70	125.67	0.0194	0.1176	62
Grok-4.1 Fast	48.33	297.94	197.54	0.1460	0.0925	54
Kimi-K2 (Thinking)	58.33	216.51	184.01	0.0964	0.1255	71
OpenAI-5.1 Mini	53.33	93.04	92.11	0.1062	0.1331	59
Qwen-235B (Thinking)	30.00	371.90	197.53	0.1919	0.1551	43
GPT-5.2	64.00	283.88	260.88	0.1774	0.0887	64
Average (8 models)	50.96	263.34	185.40	0.1055	0.1209	58.63
<i>Framework: Reflection (Step-Level)</i>						
DeepSeek-V3.2 (Exp.)	–	–	–	–	–	–
Gemini-3 (Fast)	–	–	–	–	–	–
GLM-4.6	51.67	92.35	77.18	0.0148	0.1072	57
Grok-4.1 Fast	–	–	–	–	–	–
Kimi-K2 (Thinking)	51.67	181.19	111.29	0.0353	0.1362	53
OpenAI-5.1 Mini	–	–	–	–	–	–
Qwen-235B (Thinking)	–	–	–	–	–	–
GPT-5.2	56.00	398.71	324.01	0.1536	0.1048	56
Average (8 models)	53.11	224.08	170.83	0.0679	0.1161	55.33
<i>Framework: Plan-and-Act</i>						
DeepSeek-V3.2 (Exp.)	–	–	–	–	–	–
Gemini-3 (Fast)	–	–	–	–	–	–
GLM-4.6	48.33	231.35	113.01	0.0198	0.1096	55
Grok-4.1 Fast	–	–	–	–	–	–
Kimi-K2 (Thinking)	48.67	170.26	105.36	0.0000	0.1056	51
OpenAI-5.1 Mini	–	–	–	–	–	–
Qwen-235B (Thinking)	–	–	–	–	–	–
GPT-5.2	64.00	323.88	193.02	0.0152	0.1096	64
Average (8 models)	53.67	241.83	137.13	0.0117	0.1083	56.67
<i>Heuristic Policy (Upper Bound, Easy)</i>						
Hand-crafted Policy	180.00	674.18	729.46	0.0266	0.007	180

Table 8: Performance comparison of eight large language models under four agent frameworks in the EASY environment. A hand-crafted heuristic policy is included as an approximate upper bound.